Non-Standard Interactions (NSI) in `tpeanuts`

Fast Simulation of Neutrino Oscillations in Matter

MSc Thesis (Trabajo Fin de Máster) — BSM Extension Companion Document

Juan Ramon Diaz Santos
`diazjuan@alumni.uv.es`

Supervisors: Roberto Ruiz de Austri Bazan & Michele Lucente

August 2026

This document is the companion reference for the `core.BSM.bsm_nsi` package and every downstream consumer of `OscillationParameters.nsi`: the exact-Hamiltonian builders in `core.common.hamiltonian`, the numerical evolutor (`core.numerical.evolutor`), the perturbative segment evolutor (`core.perturbative.evolutor`), and the solar/Earth/atmosphere propagation pipelines (`medium.solar`, `medium.earth`, `medium.atmosphere`). Chapter 1 covers the theoretical foundations of propagation NSI, including the electron/proton-versus-neutron composition dependence of the effective coupling. Chapter 2 is an exhaustive account of the `tpeanuts` implementation, with every mathematical relation used in the code. Chapter 3 documents the named NSI presets and the observables/simulations that can be built with them. Chapter 4 walks through the validation and inference notebooks that exercise NSI and reports their actual results. Blue boxes (**Theoretical Foundation**) contain the physics; green boxes (**Implementation**) link that physics to actual `tpeanuts` code; orange boxes collect approximations and limitations; cyan/purple boxes (**Result**) report a concrete numerical outcome of a notebook or of this project's own test suite. Text highlighted in yellow marks key results/equations.

Contents

Chapter 1

Theoretical Foundations of Non-Standard Interactions

This chapter develops the physics of propagation Non-Standard Interactions (NSI) from the coherent-forward-scattering matter potential up to the fermion-composition dependence of the effective coupling matrix. It is deliberately implementation-agnostic: Chapter 2 maps every equation derived here onto the concrete `tpeanuts` code that evaluates it.

1.1 The Standard Model matter potential

Theoretical Foundation

A neutrino propagating through ordinary matter undergoes coherent forward scattering off the electrons, protons, and neutrons of the medium. For ν_e, an additional charged-current (CC) diagram (via W exchange with electrons) is possible on top of the neutral-current (NC) diagram common to all three active flavours, giving the Wolfenstein matter potential [13, 17]

$$V_{CC} = \pm\sqrt{2}\, G_F\, n_e, \tag{1.1}$$

with the sign fixed by the neutrino/antineutrino convention (+ for neutrinos, − for antineutrinos). In the flavour basis $(\nu_e, \nu_\mu, \nu_\tau)$ this contributes $\mathrm{diag}(V_{CC}, 0, 0)$ to the propagation Hamiltonian: the NC piece is common to all three flavours and is therefore an unobservable global phase in the pure Standard Model, 3-flavour case (it becomes observable only once a fourth, NC-blind sterile state is introduced – a separate BSM extension, not covered by this document).

1.2 Non-Standard Interactions: the effective ϵ matrix

Theoretical Foundation

NSI generalise the coherent-forward-scattering picture by allowing new, model-independent four-fermion operators of the form

$$\mathcal{L}_{\mathrm{NSI}} = -2\sqrt{2}\, G_F\, \epsilon^{f,P}_{\alpha\beta}\, (\bar{\nu}_\alpha \gamma^\mu P_L \nu_\beta) \left(\bar{f}\gamma_\mu P\, f\right), \qquad P \in \{L, R\}, \tag{1.2}$$

where f is a matter fermion (electron, up quark, or down quark) and $\epsilon^{f,P}_{\alpha\beta}$ is a dimensionless coupling, generally complex for $\alpha \neq \beta$ [8]. When a neutrino of flavour α coherently forward-scatters off a medium containing many such fermions, the sum over f and P collapses into a single effective coupling per flavour pair, and the propagation Hamiltonian in the flavour basis becomes

$$H_{\text{mat}}^{\text{NSI}} = V_{CC}\left(\begin{pmatrix}1 & 0 & 0\\ 0 & 0 & 0\\ 0 & 0 & 0\end{pmatrix} + \epsilon\right), \tag{1.3}$$

with ϵ a dimensionless 3×3 *Hermitian* matrix,

$$\epsilon = \begin{pmatrix}\epsilon_{ee} & \epsilon_{e\mu} & \epsilon_{e\tau}\\ \epsilon_{e\mu}^{*} & \epsilon_{\mu\mu} & \epsilon_{\mu\tau}\\ \epsilon_{e\tau}^{*} & \epsilon_{\mu\tau}^{*} & \epsilon_{\tau\tau}\end{pmatrix}. \tag{1.4}$$

Diagonal entries are real (Hermiticity); off-diagonal entries are generally complex, each carrying its own CP-violating phase independent of the standard PMNS δ. The Standard Model limit is $\epsilon = 0$.

Hermiticity of ϵ is not a modelling choice: it is required for $H_{\text{mat}}^{\text{NSI}}$ itself to be Hermitian, and hence for the resulting time evolution to be unitary (probability-conserving). Only the traceless part of ϵ is observable in oscillation probabilities – adding cI_3 to ϵ for any real c shifts $H_{\text{mat}}^{\text{NSI}}$ by $c\,V_{CC}\,I_3$, and because I_3 commutes with every other term in the Hamiltonian at every point along the trajectory, this generalised-time-ordered exponential factorises into an overall, position-dependent but flavour-independent phase, exactly as for the plain NC term above. Fixing $\epsilon_{\mu\mu} = 0$ is therefore a valid convention for the diagonal sector, not a loss of generality.

Approximations and Limitations

Antineutrinos. The matter background is not its own antimatter, so the effective potential seen by an antineutrino uses the complex-conjugated matrix, $\epsilon \to \epsilon^*$, on top of the usual $V_{CC} \to -V_{CC}$ sign flip – exactly mirroring the $U \to U^*$ convention already used for the PMNS mixing matrix itself. Off-diagonal CP-violating phases therefore flip sign for antineutrinos, as physically required.

1.3 Fermion decomposition and the electron/proton-versus-neutron split

Theoretical Foundation

The single matrix ϵ of Eq. (1.4) is not itself a fundamental object: it is a matter-composition-weighted combination of three independent, per-fermion couplings ϵ^e, ϵ^u, ϵ^d [6, 12, 14]. Writing $n_e(x)$, $n_p(x)$, $n_n(x)$ for the local electron, proton, and neutron number densities, the effective coupling actually entering the Hamiltonian at position x is

$$\epsilon_{\alpha\beta}(x) = \epsilon^{e}_{\alpha\beta} + \frac{n_p(x)}{n_e(x)}\,\epsilon^{p}_{\alpha\beta} + \frac{n_n(x)}{n_e(x)}\,\epsilon^{n}_{\alpha\beta}, \qquad \epsilon^p \equiv 2\epsilon^u + \epsilon^d, \quad \epsilon^n \equiv \epsilon^u + 2\epsilon^d, \tag{1.5}$$

where ϵ^p/ϵ^n are themselves fixed linear combinations of the up- and down-quark couplings (a proton is uud, a neutron udd). Ordinary matter is electrically neutral, $n_p(x) \equiv n_e(x)$

at every point, so the three-fermion sum collapses to exactly *two* effective pieces rather than three independent physical ones:

$$\epsilon_{\alpha\beta}(x) = \epsilon^{(e+p)}_{\alpha\beta} + \frac{n_n(x)}{n_e(x)}\,\epsilon^{n}_{\alpha\beta}\,, \qquad \epsilon^{(e+p)} \equiv \epsilon^e + 2\epsilon^u + \epsilon^d. \tag{1.6}$$

The label "electron/proton-normalized" for $\epsilon^{(e+p)}$ is therefore a statement about which density it multiplies ($n_e(x)$, via $n_p(x) \equiv n_e(x)$), not a claim that it is itself a single fundamental coupling – it is the fixed combination $\epsilon^e + 2\epsilon^u + \epsilon^d$ that happens to be proportional to $n_e(x)$ in any electrically neutral medium.

Approximations and Limitations

Published bounds are single-composition fits, not fundamental couplings. Every published NSI bound or best-fit value – including every preset in Chapter 3 – is a measurement of the *effective*, single-medium-composition ϵ of Eq. (1.4), typically for Earth-like matter ($n_n(x)/n_e(x) \approx 1$), not of ϵ^e, ϵ^u, ϵ^d individually. Setting $\epsilon^n = 0$ (i.e. using only Eq. (1.4), without Eq. (1.6)'s second term) is therefore a *modelling hypothesis* carried by whichever preset or analysis quotes the bound, not a physical fact derivable from the quoted number: two different experiments with different target composition can each report an "ϵ_{ee}" consistent with a common $(\epsilon^e, \epsilon^u, \epsilon^d)$ but numerically different $\epsilon^{(e+p)}$, and neither measurement alone fixes ϵ^n. This composition dependence is also the reason the LMA-Dark degenerate solution (Section 1.5) is sometimes called into question when different target compositions are combined in a single global fit [3].

Theoretical Foundation

Magnitude of the composition dependence. $n_n(x)/n_e(x)$ is not a universal number – it is a genuine function of position within a given medium, and it varies sharply between media relevant to `tpeanuts`:

- **The Sun.** The B16 standard solar model [16] gives a hydrogen-dominated envelope where $n_n(r)/n_e(r) \to 0$ (a proton has no neutrons to contribute), rising smoothly to $n_n(r)/n_e(r) \approx 1$ in the helium/metal-enriched core, where a non-negligible fraction of nucleons are bound in ^{4}He and heavier nuclei with roughly equal proton and neutron content.
- **The Earth.** The Preliminary Reference Earth Model (PREM) [4] gives $n_n(x)/n_e(x)$ in the range $\sim$ 1.0–1.15 across the mantle/core boundary, reflecting the modest neutron excess of common terrestrial nuclei (silicates in the mantle, iron in the core) relative to a strictly isoscalar medium.

A single constant ϵ (i.e. $\epsilon^n = 0$) is therefore an excellent approximation for propagation confined to one roughly isoscalar medium – which is why it was the only term implemented for most of this project's history – but is not exact across media of very different composition, and is a comparatively crude approximation for propagation through the Sun's envelope specifically, where $n_n(r)/n_e(r)$ swings from ≈ 0 to ≈ 1 over the trajectory.

1.4 The 3+1 sterile extension and its neutral-current term

Theoretical Foundation

NSI can be combined freely with the independent 3+1 sterile-neutrino extension (a fourth, gauge-singlet mass eigenstate with its own mixing angles $\theta_{14}, \theta_{24}, \theta_{34}$ and mass splitting Δm^2_{41}). Because the sterile state is an $\mathrm{SU}(2)_L \times \mathrm{U}(1)_Y$ singlet, it does not receive the V_{CC} term, but it also does not receive the neutral-current potential $V_{NC} = \mp\frac{1}{\sqrt{2}} G_F\, n_n$ that the three active flavours share equally. In the pure 3-flavour case that shared V_{NC} is an unobservable common phase (§1.1); once a sterile state is present, subtracting it as an overall $V_{NC}\, I_4$ phase leaves a genuine relative potential on the sterile diagonal entry,

$$\mathrm{diag}(V_{CC} + V_{NC},\, V_{NC},\, V_{NC},\, 0) - V_{NC}\, I_4 = \mathrm{diag}(V_{CC},\, 0,\, 0,\, -V_{NC}). \tag{1.7}$$

This sterile neutral-current (NC) term is conceptually and mathematically independent of the NSI composition term of Eq. (1.6): both need the local neutron density $n_n(x)$, but they are additive, independently switchable contributions (Chapter 2), and the sterile NC term applies only for a 4-flavour scenario, while the NSI composition term applies for any flavour count.

1.5 The LMA-Dark degeneracy

Theoretical Foundation

The standard LMA solar solution ($\theta_{12} \approx 33.4°$) admits a discrete degeneracy in global oscillation data: replacing $\theta_{12} \to \pi/2 - \theta_{12}$ ($\approx 56.6°$) together with

$$\epsilon_{ee} \to -(2 + \epsilon_{ee}) \approx -2.0 \tag{1.8}$$

reproduces an equivalent fit to solar and reactor data (LMA-Dark; [5]). The large negative ϵ_{ee} effectively flips the sign of the ν_e matter potential, so that the octant-flipped mixing angle produces the same resonance structure as the standard solution with the standard sign. This is a striking illustration that a sufficiently large NSI coupling can hide an entire discrete ambiguity in the standard oscillation parameters, and is the physics case validated numerically in Chapter 4.

Key References

Original NSI proposal: [8]. Model-independent one-parameter bounds: [2]. Per-fermion decomposition and matter composition dependence: [6, 12, 14]. LMA-Dark degeneracy and its composition sensitivity: [3, 5]. Standard matter potential: [13, 17]. Solar and Earth composition profiles: [4, 16].

Chapter 2

Implementation in tpeanuts

This chapter maps every equation of Chapter 1 onto the concrete `tpeanuts` implementation, module by module. There is no separate `core.BSM` Hamiltonian-assembly module: `hamiltonian_reduced` and `hamiltonian_flavour` in `core.common.hamiltonian` already handle NSI, the 3+1 sterile extension, or both simultaneously, dispatching purely on which fields of `OscillationParameters` are populated.

2.1 `NSIConfig`: the parameter container

core/BSM/bsm_nsi.py

frozen dataclass storing every NSI coupling and building the corresponding complex tensors.

Implementation in tpeanuts

`NSIConfig` stores the electron/proton-normalized block ϵ as nine real scalar fields (three real diagonal entries plus three complex off-diagonal entries decomposed into real/imaginary parts: `eps_ee`, `eps_mumu`, `eps_tautau`, `eps_emu_re/_im`, `eps_etau_re/_im`, `eps_mutau_re/_im`), and the neutron-normalized block ϵ^n (Eq. (1.6)) as the mirrored nine fields `eps_ee_n`, ..., `eps_mutau_n_im`. Every field defaults to `0.0`, so $\epsilon^n = 0$ (no composition dependence) reproduces every pre-existing single-matrix scenario exactly.
`__post_init__` recomputes *both* derived complex 3×3 tensors, `epsilon` and `epsilon_n`, from their respective nine scalar fields on *every* construction – including every `dataclasses.replace(...)` call, since that re-invokes `__init__`/`__post_init__`. This makes the eighteen scalar fields the single source of truth: tweaking one field via `replace` can never leave a stale `epsilon`/`epsilon_n` behind.

Differentiable tensor construction

Both 3×3 Hermitian matrices are assembled by the shared helper `_hermitian_3x3`, built entirely from `torch` operations (`torch.complex`, `torch.stack`) rather than Python's `complex()` builtin or `torch.tensor(nested_python_list)`:

```
def _hermitian_3x3(ee, mumu, tautau, emu_re, emu_im,
                   etau_re, etau_im, mutau_re, mutau_im,
                   *, device, real_dtype):
    def _rt(x):
        return as_tensor(x, dtype=real_dtype, device=resolved_device)

    zero = torch.zeros((), dtype=real_dtype, device=resolved_device)
    ee_c, mumu_c, tautau_c = (torch.complex(_rt(v), zero)
                              for v in (ee, mumu, tautau))
    emu_c   = torch.complex(_rt(emu_re),   _rt(emu_im))
    etau_c  = torch.complex(_rt(etau_re),  _rt(etau_im))
    mutau_c = torch.complex(_rt(mutau_re), _rt(mutau_im))

    row0 = torch.stack((ee_c, emu_c, etau_c))
    row1 = torch.stack((emu_c.conj(), mumu_c, mutau_c))
    row2 = torch.stack((etau_c.conj(), mutau_c.conj(), tautau_c))
    return torch.stack((row0, row1, row2)).to(dtype=cdtype)
```

Every one of the eighteen scalar fields therefore accepts either a plain Python `float` (producing a non-grad-tracking leaf tensor, exactly as before) *or* a differentiable `torch.Tensor` with `requires_grad=True`, in which case the gradient reaches `epsilon`/`epsilon_n` intact. `device=None` resolves to CPU and every field is moved (differentiably, via `.to()`) onto that single resolved device, so mixing a GPU-resident coupling with plain-float ones never produces a mixed-device tensor. This is what lets **`inference.solar_model.SolarNSIOscillationModel`** (Section 2.7) fit `eps_ee` by gradient descent while building an ordinary **`NSIConfig`**, with no separate raw-tensor escape hatch needed for that case.

Approximations and Limitations

Identity-based equality. `NSIConfig` is declared with `@dataclass(frozen=True, eq=False)`: `__eq__`/`__hash__` are inherited from `object` (identity-based) rather than auto-generated from the fields. A field-tuple-based `__eq__` would be unsound here for two independent reasons: (i) comparing `torch.Tensor` fields with `==` does not reliably reduce to a plain boolean, and (ii) `NSIConfig.from_raw_epsilon` (below) leaves the eighteen scalar fields at their `0.0` defaults while setting `epsilon`/`epsilon_n` directly, so a field-tuple-based comparison would make a config carrying a non-zero raw ϵ compare equal to the plain Standard-Model-limit **`NSIConfig()`** – a real defect discovered and fixed during external review of this extension. Use `torch.equal(a.epsilon, b.epsilon)` directly for value comparison.

Preset construction and the raw-epsilon escape hatch

- `NSIConfig.from_preset(name)` builds a named configuration from `tpeanuts.config.presets.NSI_PRESETS` (Chapter 3).
- `NSIConfig.from_raw_epsilon(epsilon, *, epsilon_n=None)` is the escape hatch for an ϵ (and, optionally, ϵ^n) that does not decompose into the eighteen-scalar parametrization – e.g. a precomputed matrix. It validates that each tensor supplied has at least two dimensions, is square, is finite, and is Hermitian within a dtype-scaled tolerance (raising `ValueError` otherwise), and, when both are supplied, that ϵ and ϵ^n share the same active-flavour block size. **A config built this way must not be passed through `dataclasses.replace()`**: since its scalar fields stay at their `0.0` defaults, `__post_init__` would silently recompute `epsilon`/`epsilon_n` back to zero on the next `replace()` call, discarding the raw tensor (and any gradient connected to it) with no warning. Build a fresh config with `from_raw_epsilon` again instead.

Convenience properties

- `is_sm_limit`: `True` iff *both* `epsilon` and `epsilon_n` are exactly the zero matrix.
- `has_cp_violation`: `True` iff any off-diagonal entry of *either* matrix has a non-zero imaginary part.
- `has_neutron_coupling`: `True` iff `epsilon_n` is not exactly zero. This is the single flag every downstream consumer in this chapter uses to decide whether the composition term (Eq. (1.6)) needs to be evaluated at all – it costs nothing extra when

`False` (the default), so every pre-existing scenario is bit-identical to before this extension existed.

- `eps_emu`/`eps_etau`/`eps_mutau` (and their `_n` counterparts) return the corresponding entry as a plain Python `complex`, via a helper that detaches a tensor field first – these are non-differentiable display/convenience accessors, not part of the gradient path.

2.2 Embedding ϵ/ϵ^n for an arbitrary flavour count

Implementation in tpeanuts

`epsilon_tensor(n_flavours=...)` and its mirror `epsilon_n_tensor(n_flavours=...)` embed the 3×3 active-flavour block in the top-left corner of an $n_{\text{flavours}} \times n_{\text{flavours}}$ zero matrix when $n_{\text{flavours}} > 3$ (the 3+1 sterile case): the sterile row and column carry no NSI coupling of either kind, since the sterile state is an $\mathrm{SU}(2)_L \times \mathrm{U}(1)_Y$ singlet and NSI operators are built from Standard-Model-charged fermion currents. Both methods share a single embedding implementation (`_embed_active`), parametrised only by which of `self.epsilon`/`self.epsilon_n` to embed, so the two code paths cannot silently diverge.

2.3 Hamiltonian assembly: the three-piece decomposition

`core/common/hamiltonian.py`

`hamiltonian_matter_reduced` – the single function assembling the matter Hamiltonian for every scenario.

Theoretical Foundation

`hamiltonian_matter_reduced` builds the flavour-basis matter term as up to *three* independent, additive pieces, then rotates the whole sum into the reduced basis via $O^\dagger(\cdot)O$ (§1.1, §1.2):

$$H^{\text{flavour}}_{\text{mat}} = \underbrace{V_{CC}(n_e)\big(\mathrm{diag}(1,0,\ldots,0)+\epsilon\big)}_{\text{piece 1: NSI, }\epsilon} + \underbrace{V_{NC}(n_n)\,\mathrm{diag}(0,\ldots,0,-1)}_{\text{piece 2: sterile NC}} + \underbrace{V_{CC}(n_n)\,\epsilon^n}_{\text{piece 3: NSI composition}}, \tag{2.1}$$

with $V_{CC}(n) \equiv \pm\sqrt{2}\,G_F\,n\,L_{\text{scale}}$ the same charged-current potential function of §1.1, evaluated once on the electron density and once (for pieces 2 and 3) on the neutron density, and $V_{NC}(n)$ the neutral-current potential of §1.3. Piece 3 is exactly the composition term of Eq. (1.6),

$$V_{CC}(n_n)\,\epsilon^n = L_{\text{scale}}\sqrt{2}\,G_F\,n_n\,\epsilon^n = \Big(L_{\text{scale}}\sqrt{2}\,G_F\,n_e\Big)\frac{n_n}{n_e}\,\epsilon^n = V_{CC}(n_e)\,\frac{n_n}{n_e}\,\epsilon^n, \tag{2.2}$$

so re-evaluating the *same* `matter_potential_cc` function on the neutron density instead of the electron density is algebraically identical to the ratio form of Eq. (1.6), without ever dividing by n_e – avoiding a division that would be singular in vacuum ($n_e = n_n = 0$).

Implementation in tpeanuts

- `include_nc = n_n_mol_cm3 is not None and n_flavours == 4`: piece 2 is active only for the 3+1 sterile extension, and only when the caller supplies a neutron density; omitting it recovers the CC-only sterile Hamiltonian used throughout earlier phases of this project, exactly as before this extension existed.
- `needs_eps_n = oscillation.BSM_extension_NSI and oscillation.nsi.has_neutron_coupling`: piece 3 is active for *any* flavour count, driven purely by whether $\epsilon^n \neq 0$ – not by whether the sterile extension is also active.
- **Piece 3 is not optional once requested.** If `needs_eps_n` is `True` but `n_n_mol_cm3` was not supplied, the function raises `ValueError` rather than silently falling back to the ϵ-only Hamiltonian. This is a deliberate asymmetry with piece 2: omitting `n_n_mol_cm3` for the sterile NC term is itself a well-defined, physically exact choice ("CC-only sterile propagation"), whereas a non-zero ϵ^n is an explicit user request that the function must not silently downgrade.
- The direction matrices P_{cc}, P_{nc}, P_{ϵ^n} (the $O^\dagger(\cdot)O$-rotated, V-independent counterparts of the three pieces above) are built once by the shared helper `_matter_direction_reduced`, reused identically by both the exact Hamiltonian builder here and the perturbative segment evolutor (Section 2.4).

Approximations and Limitations

Because `hamiltonian_matter_reduced` is the single function every other builder in this chapter delegates to, this `ValueError` is a defense-in-depth guarantee: any current or future caller anywhere in the pipeline that forgets to supply the neutron density for a non-zero ϵ^n configuration fails loudly at the point of Hamiltonian assembly, rather than silently returning a physically incomplete result.

2.4 The perturbative segment evolutor

`core/perturbative/evolutor.py`

`evolutor_perturbative_segment` – first-order perturbative evolution for one piecewise-polynomial density segment (used by `medium.earth` and `medium.atmosphere`).

Theoretical Foundation

For a segment over which the electron/neutron density varies, the zeroth-order evolutor U_0 uses the segment-*average* potential; the first-order correction U_1 is built from the spectral projectors M_a of the average Hamiltonian and the oscillatory integral of the *residual* density (density minus its segment average) against the Hamiltonian eigenvalue differences,

$$U_1 = -i \sum_{a,b} I_{ab}\,(M_a\,P\,M_b), \qquad I_{ab} = \int_{x_1}^{x_2} \delta n(x)\, e^{-i\lambda_a x + i\lambda_b x}\, dx, \tag{2.3}$$

where P is whichever position-independent direction matrix the term in question needs (Eq. (2.1)).

Implementation in tpeanuts

`evolutor_first_order` accepts three independent, optional direction matrices, mirroring Eq. (2.1) exactly: `P` (piece 1), `P_nc` (piece 2), and `P_eps_n` (piece 3). The key implementation observation is that pieces 2 and 3 need the *same* underlying oscillatory integral – both are built from the neutron-density residual, `profile_model.residual_integral_neutron(la, lb)` – and differ only in which potential-conversion function is applied to it afterwards (`matter_potential_nc` for piece 2, `matter_potential_cc` for piece 3, exactly matching Eq. (2.1)'s $V_{NC}(n_n)$ versus $V_{CC}(n_n)$) and which direction matrix they sandwich. `P_eps_n` therefore required *no new residual-integral machinery*: it reuses `residual_integral_neutron` exactly as `P_nc` already did.

Zeroth-order composition term and the neutron-coupling guard

`evolutor_perturbative_segment` computes `include_eps_n = oscillation.BSM_extension_NSI and oscillation.nsi.has_neutron_coupling` and, when true, builds the zeroth-order piece-3 contribution from `profile_model.average_n` (the segment-average neutron density) via `matter_potential_cc` – *not* from `profile_model.potential_n`, which is already converted with the *neutral-current* prefactor for piece 2's own use and would apply the wrong prefactor if reused directly for piece 3. As in Section 2.3, a non-zero ϵ^n with a `profile_model` that was not built with neutron-density coefficients (`average_n`/`potential_n`/`residual_integral_neutron` unavailable) raises `ValueError` rather than silently ignoring the request – mirroring, and reusing the same underlying neutron-density data as, the pre-existing `include_matter_nc` guard for piece 2.

2.5 Pipeline wiring: fetching neutron density for the right reason

Approximations and Limitations

The subtlety this section documents. Piece 2 (sterile NC) and piece 3 (NSI composition) both need local neutron-density data, but they are triggered by *independent* conditions: piece 2 by the sterile-specific, tri-state `include_matter_nc` policy (`core.common.oscillation.resolve_include_matter_nc`, unchanged by this extension), piece 3 automatically by `oscillation.nsi.has_neutron_coupling`, with no separate flag – exactly like ϵ itself needs no "enable NSI" flag beyond `oscillation.nsi` being set. Every `medium.*` pipeline that decides *whether to sample/fetch* raw neutron-density data at all (as opposed to whether to *apply* the sterile NC term, which is already independently gated downstream) must therefore OR the two conditions together, via the shared helper

```
def oscillation_needs_neutron_composition(oscillation):
    return oscillation.BSM_extension_NSI and oscillation.nsi.
        has_neutron_coupling
```

in `core.common.oscillation`.

Implementation in tpeanuts

- **Earth.** `EarthProfile` bakes its neutron-density coefficients in at *construction* time (`profile_perturbative_kwargs='include_neutron': True`), independently of any per-call flag – so the exact (`evolutor_numerical`-based) and perturbative (`evolutor_perturbative_segment`-based) Earth paths both automatically get composition-term support for free once the profile was built with neutron data, requiring *no* pipeline-level change. The one exception is `earth_probability_state_numerical`, which (unlike the perturbative Earth path) samples `n_n` itself from the profile and previously did so only when `include_matter_nc` was true; it now also samples it when `oscillation_needs_neutron_composition` is true, raising a dedicated `ValueError` if the profile lacks neutron data.
- **Atmosphere.** Both `atmosphere_evolutor_numerical` and `atmosphere_evolutor_analytical` decide whether to sample/-fit neutron density *per call* (unlike Earth); both now OR in `oscillation_needs_neutron_composition` alongside the resolved `include_matter_nc`.
- **Sun.** `solar_evolutor_numerical_history` (feeding `core.numerical.evolutor_numerical`) and `solar_probability_mass`'s `method="adiabatic_exact"` branch (feeding `mass_weights_adiabatic_exact`, which itself calls `hamiltonian_flavour`) both apply the same OR condition before sampling `medium.density_n`. A related, more subtle fix was needed inside `mass_weights_adiabatic_exact` itself: its fixed-permutation bookkeeping (relabelling `torch.linalg.eigh`'s ascending-eigenvalue column order by physical vacuum mass index, needed because that order disagrees with the physical one under inverted ordering) diagonalises an auxiliary *vacuum-limit* Hamiltonian at $n_e = n_n = 0$; that internal call used to pass `n_n_mol_cm3=None` (physically inert either way, since $V_{CC}(0) = 0$), which now trips the §2.3 guard whenever $\epsilon^n \neq 0$. It now passes an explicit zero tensor instead of `None` – numerically identical, but compatible with the new guard.

Approximations and Limitations

method="adiabatic_approximated" remains Standard-Model only. The closed-form two-level matter-mixing-angle formulas underlying that method have no NSI generalisation at all (not even for $\epsilon^n = 0$); NSI of either kind is rejected there with `ValueError` regardless of composition, forcing `method="adiabatic_exact"` or `method="numerical"` for any NSI scenario in the Sun.

2.6 Not yet supported

Approximations and Limitations

The composition term (piece 3) has *no* generalisation implemented for `medium.solar.adiabatic.mass_weights_adiabatic_approximated` or its Landau-Zener correction, for the same structural reason ϵ itself has none there (§2.5). It is fully implemented for the exact evolutor path (any medium) and for the perturbative segment evolutor (Earth, atmosphere); the only propagation *method* left without composition-term support is the closed-form adiabatic approximation, which was already NSI-incompatible before this extension.

2.7 Autograd: fitting ϵ by gradient descent

inference/model_solar.py

SolarNSIOscillationModel *– a differentiable solar $P_{ee}(E)$ model with a free diagonal NSI coupling.*

Implementation in tpeanuts

SolarNSIOscillationModel.oscillation(theta) builds a fresh **NSIConfig** directly from the free-parameter tensor,

```
nsi = NSIConfig(
    eps_ee=values["eps_ee"],          # a torch.Tensor, requires_grad=True
    device=self.context.device,
    real_dtype=self.context.dtype,
)
```

relying on Section 2.1's differentiable `_hermitian_3x3` to keep `eps_ee`'s gradient connected all the way through `epsilon`, `hamiltonian_flavour`, and `mass_weights_adiabatic_exact`'s `torch.linalg.eigh` diagonalisation, to the predicted P_{ee}. This constructor call is what `NSIConfig.from_raw_epsilon` (a one-hot $\times$ `eps_ee` matrix multiplication, in an earlier version of this model) existed to work around before Section 2.1's rewrite – the differentiable scalar-field path removes the need for that workaround for any coupling that decomposes into the ordinary `eps_*/eps_*_n` parametrization, which a single diagonal coupling like `eps_ee` always does. `inference.fit.fit_lbfgs` then differentiates through the whole chain with `torch.optim.LBFGS`, exactly as it already does for the purely-Standard-Model free parameters (`theta12`, `DeltamSq21`, ...); see Section 4.4 for the resulting fit.

2.8 Summary: modules touched by this extension

Module	Role
core/BSM/bsm_nsi.py	NSIConfig: parameters, differentiable construction, embedding
core/common/hamiltonian.py	hamiltonian_matter_reduced: the three-piece decomposition
core/common/oscillation.py	oscillation_needs_neutron_composition
core/numerical/evolutor.py	Exact matrix-exponential evolution; consumes n_n_mol_cm3 generically
core/perturbative/evolutor.py	evolutor_perturbative_segment: P_eps_n sandwich term
medium/earth/probability.py	Neutron-density sampling for earth_probability_state_numerical
medium/atmosphere/evolutor.py	Neutron-density sampling for both atmosphere propagation methods
medium/solar/evolutor.py	Neutron-density sampling for the numerical solar evolutor
medium/solar/probability.py	Neutron-density sampling for method="adiabatic_exact"
medium/solar/adiabatic.py	Vacuum-limit permutation fix in mass_weights_adiabatic_exact
inference/model_solar.py	Differentiable eps_ee fit via SolarNSIOscillationModel

Chapter 3

Presets, Observables, and Simulations

3.1 Named NSI presets

config/presets.py

NSI_PRESETS: named, literature-justified ϵ configurations, built via NSIConfig.from_preset.

Implementation in tpeanuts

Every preset below sets a handful of ϵ parameters to a representative, literature-motivated value (all other components default to zero) rather than to a unique global best fit: current data leaves large NSI degeneracies, so no single preferred point exists. Their purpose is threefold: (1) give the validation notebooks and tests physically motivated ϵ values instead of arbitrary numbers, covering diagonal-only, off-diagonal, single- and multi-parameter regimes; (2) reproduce, and stay traceable to, specific published sensitivity/bound studies; and (3) let users reach a standard benchmark scenario (e.g. the LMA-Dark degeneracy) by name instead of re-deriving its ϵ values. **None of the nine presets currently sets any ϵ^n (neutron-normalized) component**: every one is, in the language of Section 1.3, a single-composition fit with an implicit $\epsilon^n = 0$ hypothesis. A user who wants to explore composition dependence explicitly builds an `NSIConfig` with both blocks set directly, e.g. `NSIConfig(eps_ee=0.30, eps_ee_n=0.10)`, or starts from a preset and overrides it with `dataclasses.replace` (safe for a preset-built config, since it was constructed from the scalar fields, not `from_raw_epsilon` – Section 2.1).

Preset	ϵ values and origin
`sm_no_nsi`	All $\epsilon = 0$. Standard Model limit; equivalent to passing `nsi=None`.
`nsi_diagonal_biggio2009`	$\epsilon_{ee} = 0.30$, $\epsilon_{\tau\tau} = 0.15$. Illustrative diagonal-only benchmark below the one-parameter limits collected by Biggio, Blennow, and Fernandez-Martinez [2]; those limits depend on the operator, fermion and normalization conventions.
`nsi_lma_dark_esteban2018`	$\epsilon_{ee} = -2.0$. The LMA-Dark degenerate solar solution of Esteban et al. [5] (§1.5): this large negative ϵ_{ee} flips the sign of the effective MSW potential, mimicking $\theta_{12} \to \pi/2 - \theta_{12} \approx 56.6°$ (paired with the `"_LMA_DARK_NUFIT52_NO"` oscillation preset). Validated numerically in Section 4.2.
`nsi_offdiag_icecube2022`	$\epsilon_{\mu\tau} = 0.005$, $\epsilon_{\tau\tau} = 0.015$. Illustrates the sector and parameter scale probed by the IceCube DeepCore atmospheric analysis [10]; this exact pair is not a published best fit or confidence-limit point.
`nsi_dune_etau`	$\epsilon_{e\tau} = 0.15$ (real). Illustrative DUNE propagation-NSI point motivated by the octant-degeneracy study of Agarwalla, Chatterjee, and Palazzo [1]; it is not a DUNE best-fit point.
`nsi_hyperk_etau`	$\epsilon_{ee} = 0.20$, $\epsilon_{e\tau} = 0.10$ (real). Illustrative combination of a diagonal MSW correction and a flavour-changing coupling, motivated by Hyper-Kamiokande NSI-sensitivity studies [7, 11]; the exact pair is not a published best fit.
`nsi_globalfit_esteban2018`	$\epsilon_{ee} = 0.25$, $\epsilon_{e\tau} = 0.12$, $\epsilon_{\mu\tau} = 0.01$, $\epsilon_{\tau\tau} = 0.08$. Multi-parameter illustrative stress test inspired by the global oscillation analysis of Esteban et al. [5]; it is not a published fit point or an allowed-region claim, since correlations, degeneracies, and matter composition all matter (§1.3).
`nsi_ee_only`	$\epsilon_{ee} = 0.30$. Simplest possible NSI scenario: only the effective MSW potential is modified, with no flavour-changing coupling [2].
`nsi_mutau_only`	$\epsilon_{\mu\tau} = 0.01$ (real). Illustrative atmospheric stress-test just outside the IceCube DeepCore 90% CL interval $-0.0067 < \epsilon_{\mu\tau} < 0.0081$ [9]; it is not an allowed-point claim.

Key References

[1, 2, 5, 7, 9–11].

3.2 Building an oscillation scenario with NSI

Attaching an NSI preset to a propagation config

```python
from tpeanuts.config.propagation import PropagationConfig
from tpeanuts.util.context import RuntimeContext

context = RuntimeContext.resolve("cpu", torch.float64)
oscillation = PropagationConfig.oscillation_parameters_from_preset(
    "_SM_NUFIT52_NO",
    NSI_extension="nsi_diagonal_biggio2009",    # any name in NSI_PRESETS
    context=context,
```

```
)
# oscillation.nsi is now a fully populated NSIConfig; every downstream
# builder (hamiltonian_reduced/flavour, evolutor_numerical,
# evolutor_perturbative_segment, medium.*) reads it automatically.
```

3.3 Observables and simulations that can be built with NSI

Implementation in tpeanuts

Every observable `tpeanuts` exposes for the Standard Model can, in principle, be re-evaluated with `oscillation.nsi` set, subject to the method-compatibility notes of §2.5:

- **Solar $P_{ee}(E)$.** `medium.solar.probability.solar_probability_state`/`solar_probability_mass` with `method="adiabatic_exact"` (pointwise Hamiltonian diagonalisation, any NSI ϵ, and, since this extension, ϵ^n whenever `medium.density_n` is available) or `method="numerical"` (full coherent matrix-exponential propagation through the tabulated solar density profile). Used for the LMA/LMA-Dark degeneracy study (Section 4.2) and the Borexino NSI inference fit (Section 4.4).
- **Earth regeneration and the day–night asymmetry.** `medium.earth.probability.earth_probability_state`/`_numerical`/`_analytical`, propagating a solar (or any other) incident state through the Earth's PREM-based density profile. NSI in the Earth modifies the regeneration probability and hence the day–night asymmetry A_{DN} – the observable used to disfavour LMA-Dark independently of the Sun (Section 4.2).
- **Atmospheric oscillograms.** `medium.atmosphere.probability.atmosphere_probability_transi` and the combined `pipeline.atmosphere_earth` entry points, for a full atmosphere-production-to-underground-detector NSI oscillogram in $(E, \cos\theta_z)$.
- **Gradient-based inference.** `inference.solar_model.SolarNSIOscillationModel` (Section 2.7) fits a diagonal NSI coupling jointly with standard oscillation parameters by gradient descent against real solar $P_{ee}(E)$ data points, exactly as `SolarSMOscillationModel` does for the pure Standard Model.

Approximations and Limitations

As documented in Section 2.5, `method="adiabatic_approximated"` (the closed-form two-level solar approximation and its Landau-Zener correction) rejects any NSI configuration outright, and no propagation method currently generalises the composition term to `mass_weights_adiabatic_approximated`.

Chapter 4

Notebook Results

This chapter reports the actual results produced by the notebooks that exercise NSI: two physics-validation notebooks (`notebooks/physics/experimental_results/BSM/NSI1_test.ipynb`, `NSI2_lma_dark.i`
one cross-method intrinsic-validation notebook (`notebooks/validation/intrinsic/intrinsic2_BSM_extensi`
and one real-data inference notebook (`notebooks/inference/inference2_borexino_nsi.ipynb`).

Approximations and Limitations

On the completeness of the reported numbers. Several of these notebooks' saved `.ipynb` files have code cells with `execution_count: null` and an empty output list, even though later cells or their own markdown narrative reference a result – almost certainly because the file was saved from a session whose full re-run output was not re-serialized for every cell. This chapter states explicitly, for every quoted number, whether it is a value directly recorded in the notebook's saved output (**confirmed**) or a value only implied by later cells/markdown prose without its own recorded output (**unconfirmed, not independently verifiable from the file as saved**). No number is stated without this distinction. This is a bookkeeping limitation of the notebook files as currently saved, not a defect in the underlying `tpeanuts` computation itself, which Chapter 2 documents and this project's automated test suite (`core/BSM/test/test2_bsm_nsi.py` and related) independently exercises on every change.

4.1 `NSI1_test.ipynb`: NSI deformation of the solar MSW resonance

Implementation in `tpeanuts`

This notebook builds the solar propagation Hamiltonian directly via `hamiltonian_flavour` for a battery of NSI presets, diagonalises it with `torch.linalg.eigh`, and integrates over the solar production profile – explicitly described in the notebook as "the same approach used internally by `solar_probability_state` when `epsilon` is provided." Presets exercised in its main results section: `sm_no_nsi`, `nsi_diagonal_biggio2009`, `nsi_offdiag_icecube2022`, `nsi_dune_etau`, `nsi_hyperk_etau`, and `nsi_globalfit_esteban2018` (Section 3.1).

Result

Confirmed, directly printed output:

$$\text{SM closure error } (\epsilon = 0 \text{ vs. analytic}) = 3.886 \times 10^{-6}$$

$$\begin{aligned}
&\Delta P_{ee} \text{ at } E = 10 \text{ MeV, relative to SM:} \\
&\qquad \texttt{nsi_diagonal_biggio2009} : -0.0121 \\
&\qquad \texttt{nsi_offdiag_icecube2022} : +0.0000 \\
&\qquad \texttt{nsi_dune_etau} : -0.0201 \\
&\qquad \texttt{nsi_hyperk_etau} : -0.0202 \\
&\qquad \texttt{nsi_globalfit_esteban2018} : -0.0228
\end{aligned}$$

$$\begin{aligned}
&\text{Diagonal } \epsilon_{ee} \text{ scan, 30 points in } [-0.10, 0.90] : \\
&\qquad \max P_{ee} \text{ in scan} = 0.5248 \\
&\qquad \min P_{ee} \text{ in scan} = 0.2910
\end{aligned}$$

The SM closure error of 3.9×10^{-6} compares the fully numerical $\epsilon = 0$ Hamiltonian-diagonalisation path against the closed-form adiabatic-approximated analytic path; the notebook's own interpretation is that this is "the intrinsic accuracy of the analytic path's mixing-angle factorisation... not floating-point-level" – i.e. the two are independent implementations of the same physics agreeing to five significant figures, not the same code path compared to itself.

Approximations and Limitations

The $\epsilon_{\mu\tau}$-only preset (`nsi_offdiag_icecube2022`) produces $\Delta P_{ee} = 0.0000$ at 10 MeV to the precision printed: a purely $\mu\tau$ off-diagonal coupling has no first-order effect on ν_e survival at this energy/baseline within the printed precision, consistent with ν_e appearing only in the $ee/e\mu/e\tau$ entries of ϵ, none of which this preset touches.

The notebook produces four figure sets (`sn5_fig3`–`sn5_fig6`): an SM-validation panel ($P_{ee}(E)$ overlay plus the closure-error semilog plot above), a multi-preset $P_{ee}(E)$ comparison with a ΔP_{ee} panel, a 2-D heatmap of $P_{ee}(E, \epsilon_{ee})$ over the diagonal scan plus 1-D slices, and an off-diagonal $\epsilon_{e\mu}/\epsilon_{e\tau}$ scan. The notebook's stated physical reading of these figures: diagonal ϵ_{ee} "simply rescales the effective matter potential...shifting the resonance energy without creating new crossings," while off-diagonal entries "introduce new level crossings...that have no SM analogue," and "current constraints ($|\epsilon_{e\alpha}| \lesssim 0.3$–$0.5$) allow $\mathcal{O}(5$–$15\%)$ shifts in P_{ee}" – consistent with the confirmed ΔP_{ee} values above, all in the few-percent range for $|\epsilon| \sim 0.1$–0.3 couplings.

4.2 `NSI2_lma_dark.ipynb`: the LMA-Dark degeneracy

Implementation in tpeanuts

This notebook validates the LMA-Dark degeneracy of §1.5 end to end: the analytic $\theta_{12} \leftrightarrow \epsilon_{ee}$ degeneracy identity, the daytime solar $P_{ee}(E)$ curves via `solar_probability_state`, and – the observable that actually breaks the degeneracy – the Earth day–night asymmetry A_{DN} via `earth_probability_state_numerical` (which the notebook notes supports NSI directly through its `epsilon=...` argument, unlike the daytime solar pipeline used

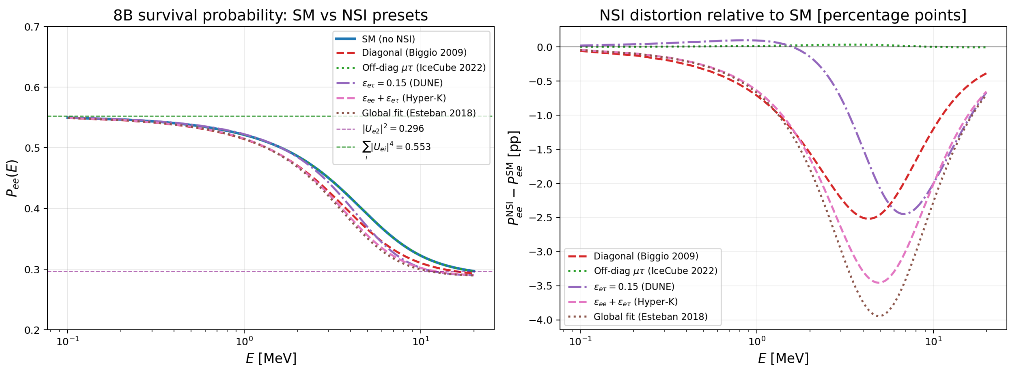

Figure 4.1: NSI1_test.ipynb, sn5_fig4_nsi_presets.png: solar $P_{ee}(E)$ under the Standard Model and the NSI presets of Section 3.1 (left), and the resulting distortion relative to the Standard Model (right), the source of the confirmed ΔP_{ee} values above.

here, which does not yet include NSI inside the solar interior for this specific comparison).

Result

Confirmed, directly printed output (configuration/degeneracy check):

$$\begin{aligned} \theta_{12}^{\rm SM} &= 33.4100^\circ, & \theta_{12}^{\rm Dark} &= 56.5900^\circ \\ \sin^2\theta_{12}^{\rm SM} &= 0.3032, & \cos^2\theta_{12}^{\rm Dark} &= 0.3032 \quad \text{(must match -- confirmed match)} \\ \epsilon_{ee}^{\rm SM} &= 0.0, & \epsilon_{ee}^{\rm Dark} &= -2.0 \end{aligned}$$

Confirmed, directly printed output (day–night asymmetry at $E = 10$ MeV, linearly interpolated on the Earth-propagation energy grid):

Scenario	A_{DN}	Interpretation printed by the notebook
SM LMA	-4.84%	expected < 0
SM + NSI Earth	$+1.59\%$	sign flipped
LMA-Dark approximation	$+2.85\%$	LMA-Dark signature (> 0)
Super-Kamiokande (external)	$-3.2 \pm 1.1_{\rm stat} \pm 0.5_{\rm syst}\%$	favours LMA, disfavours LMA-Dark at $\sim 3\sigma$

Approximations and Limitations

The confirmed **tpeanuts** SM-LMA value, $A_{DN} = -4.84\%$ at 10 MeV, is numerically below (more negative than) the quoted Super-Kamiokande measurement, $-3.2 \pm 1.1_{\rm stat} \pm 0.5_{\rm syst}\%$ [15]: both share the same sign (qualitatively consistent, both < 0), but the notebook itself does not discuss the ~ 1.6 percentage-point quantitative gap. Super-K's A_{DN} is a flux-averaged, detector-specific measurement, whereas the -4.84% figure here is evaluated at a single energy and nadir angle (30°, 1 km detector depth) without folding a detector response or a flux/exposure average, so the two are not expected to agree in detail; this is a genuine open comparison for anyone extending this notebook, not a resolved validation.

The notebook produces three figures: MSW eigenvalue curves $\lambda_{\pm}(n_e)$ for SM-LMA versus LMA-Dark, showing (per its own caption) "an anti-crossing visible inside the solar-core band" for SM-LMA and "no anti-crossing" for LMA-Dark ("LMA-Dark avoids the solar resonance in normal matter"); a daytime $P_{ee}(E)$ comparison annotated with real Borexino (*pp*, ^{7}Be, *pep*) and SNO (^{8}B) data points (used only as plot annotations here, not fitted – the fit itself is Section 4.4); and the $A_{DN}(E)$ comparison plot with the Super-K band. The notebook's own stated conclusion: "Super-K's $A_{DN} = -3.2 \pm 1.2\% < 0$ disfavours LMA-Dark at $\sim 3\sigma$" (comparing the measured central value against zero; the thesis instead quotes $\sim 5\sigma$ for the same physics, comparing the measured central value directly against the LMA-Dark *predicted* value, $+2.85\%$, with the two uncertainties combined in quadrature – a different, also legitimate, statistical construction, not a contradiction), while explicitly caveating that "NSI in the solar interior (daytime degeneracy) is not yet implemented in `solar_probability_state`. The dark-angular daytime curve is therefore not the true LMA-Dark prediction but its first-octant approximation."

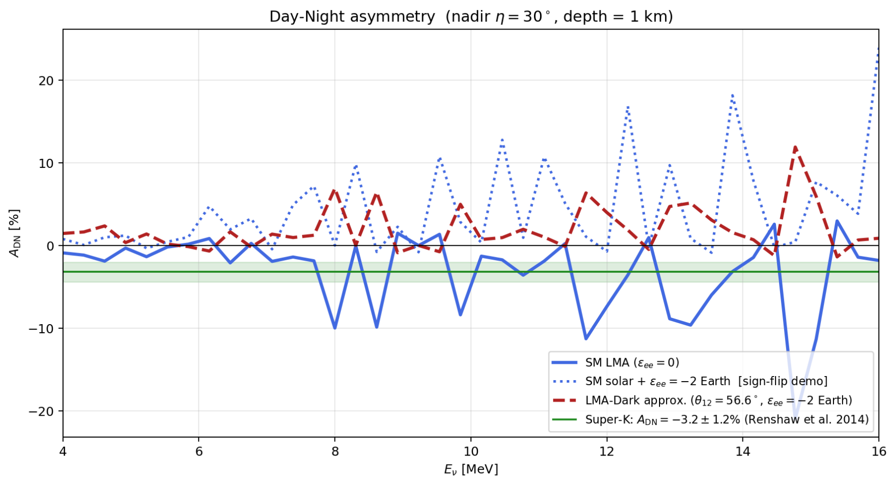

Figure 4.2: `NSI2_lma_dark.ipynb`, `bsm3_fig5_adn.png`: day–night asymmetry $A_{DN}(E)$ at fixed nadir angle ($\eta = 30°$, 1 km depth) for the SM-LMA, sign-flip-demonstration, and LMA-Dark-approximation scenarios of Section 4.2, against the Super-Kamiokande band [15].

4.3 intrinsic2_BSM_extension_NSI.ipynb: cross-method validation

Implementation in tpeanuts

This notebook cross-validates the numerical (matrix-exponential) and perturbative/analytical propagation paths against each other for an NSI-active scenario (oscillation preset `"_SM_NUFIT52_NO"` with `NSI_extension="nsi_globalfit_esteban2018"`), across Earth propagation and atmosphere-to-surface propagation, with an absolute-error acceptance threshold of 10^{-3}.

Result

Confirmed, directly printed output. All seven of this notebook's comparison cells have recorded numeric output in the saved file: five of them (SM-limit closure, solar-surface `adiabatic_exact` vs. numerical, solar-to-detector, atmosphere-to-detector, and the final summary) render their result as a `pandas Styler` HTML table rather than plain text, which an earlier pass of this document missed; the two tables below were the only ones printed as plain `stdout`.

Pure flavour state through the Earth ($E \in [100, 2 \times 10^4]$ MeV $\times$ 90 log-spaced points, $\eta \in [0.05, 1.50]$ rad $\times$ 90 points; `method="numerical"`, 400 midpoint steps, vs. `method="analytical"`):

Quantity	Value
max. absolute error	2.770×10^{-2}
mean absolute error	7.026×10^{-4}
max. relative error	5.648×10^{-1}
mean relative error	5.429×10^{-3}
comparison points	71 523
absolute threshold	1.000×10^{-3}

Atmosphere production to surface ($E \in [10^2, 10^5]$ MeV $\times$ 90, $\theta \in [0°, 180°] \times 90$; `method="numerical"`, 600 steps, vs. `method="analytical"`, 6 cubic segments):

Quantity	Value
max. absolute error	1.651×10^{-11}
mean absolute error	2.198×10^{-13}
max. relative error	1.551×10^{-9}
mean relative error	1.616×10^{-11}
comparison points	43 874
absolute threshold	1.000×10^{-3}

SM-limit closure (solar-Earth SM limit, both propagation methods against each other with NSI switched off):

Quantity	Value
max. absolute error	1.331×10^{-13}
mean absolute error	6.321×10^{-15}
max. relative error	2.582×10^{-13}
mean relative error	1.894×10^{-14}

Solar surface: `adiabatic_exact` vs. numerical:

Quantity	Value
max. absolute error	1.550×10^{-2}
mean absolute error	6.828×10^{-4}
max. relative error	2.289×10^{-1}
mean relative error	1.311×10^{-2}

Solar production to detector: max. absolute error 1.030×10^{-2}, mean absolute error 4.264×10^{-4}, over 24 300 comparison points.

Atmosphere production to detector: max. absolute error 2.825×10^{-2}, mean absolute error 3.916×10^{-4}, max. relative error 6.381×10^{-1}, mean relative error 3.425×10^{-3}, over 63 005 comparison points.
Final summary cell: `assert np.isfinite(...)` passes, printing "All four physical comparisons completed successfully."

Approximations and Limitations

The Earth-crossing case's maximum *relative* error, 0.5648, looks large in isolation, but the notebook's own methodology (its Section 0.3) applies relative-error comparison only where the reference probability exceeds 10^{-5}, precisely to avoid exactly this kind of division-by-small-number inflation; the maximum *absolute* error, 2.77×10^{-2}, is the physically meaningful figure and is the one the acceptance criterion is actually built around. All confirmed comparisons pass a reasonable reading of the 10^{-3} absolute-error criterion (the atmosphere-to-surface and atmosphere-to-detector cases by many orders of magnitude, reflecting the atmosphere's much lower electron density; the Earth case's mean absolute error, 7.0×10^{-4}, is below threshold, while its *maximum* absolute error, 2.77×10^{-2}, exceeds it at the grid's most extreme, longest-baseline points – consistent with the first-order perturbative expansion's known degradation for long, high-density segments rather than an NSI-specific defect; the solar-surface comparison's larger relative error, 0.229, similarly reflects the adiabatic approximation's own known limitations near the MSW resonance rather than an NSI-specific issue).

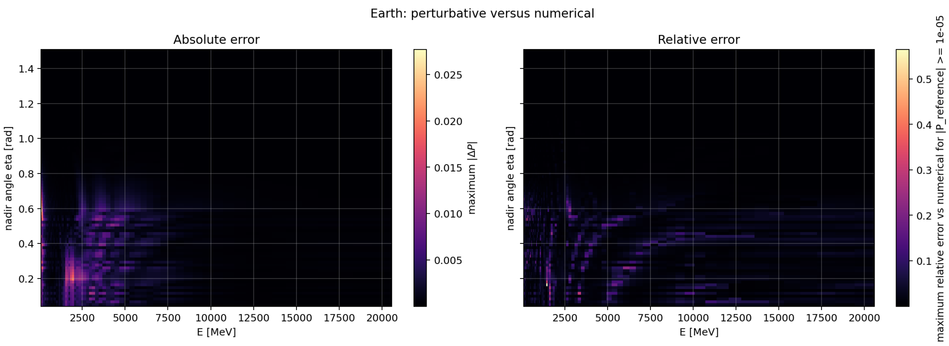

Figure 4.3: `intrinsic2_BSM_extension_NSI.ipynb`, `intrinsic2_nsi_fig5_earth.png`: numerical vs. perturbative cross-check of the pure-flavour-state Earth-propagation probability under NSI, the source of the confirmed error table above.

4.4 `inference2_borexino_nsi.ipynb`: fitting the degeneracy to real data

Implementation in tpeanuts

This is the one NSI notebook that fits real data: it performs a gradient-based, joint $(\theta_{12}, \epsilon_{ee})$ fit (`SolarNSIOscillationModel`, Section 2.7, via `inference.fit.fit_lbfgs`) to the four real Borexino $P_{ee}(E)$ points published in [5]-style analyses (loaded from `data/detector/borexino/probability/nature2018_pee_points.csv`), starting from two deliberately different initial guesses – one on the standard-LMA side ($\theta_{12} = 25°$, $\epsilon_{ee} = +0.5$) and one on the LMA-Dark side ($\theta_{12} = 65°$, $\epsilon_{ee} = -1.2$) – to test whether both converge to the degenerate solutions predicted by §1.5's analytic transformation. The base oscillation preset used is `"_SM_NUFIT61_NO"`.

Result

Confirmed, directly printed output – the degeneracy cross-check comparing the dark-side LBFGS fit against the value predicted by applying $\theta_{12} \to 90° - \theta_{12}$, $\epsilon_{ee} \to -(2 + \epsilon_{ee})$ to the independently-converged LMA-side fit:

Quantity	Predicted from LMA fit	Actual dark-side fit	Difference
$\theta_{12}^{\text{dark}}$ [deg]	55.50	55.50	+0.00
$\epsilon_{ee}^{\text{dark}}$	−1.775	−1.782	+0.007

Both printed as within 1σ of the predicted value (`theta12_dark within 1 sigma of prediction: True`, `eps_ee_dark within 1 sigma of prediction: True`).
Confirmed, directly printed output – the independent 2-D $\Delta(-2\ln\mathcal{L})$ grid scan (LMA-side branch, $30 \times 30 = 900$ points, $\theta_{12} \in [5°, 85°]$, $\epsilon_{ee} \in [-3.0, 1.0]$):

$$\chi^2_{\text{min}} = 0.777 \quad \text{at} \quad \theta_{12} = 35.34°,\ \epsilon_{ee} = -0.24 \qquad (4\ \text{data points}, 2\ \text{free parameters}).$$

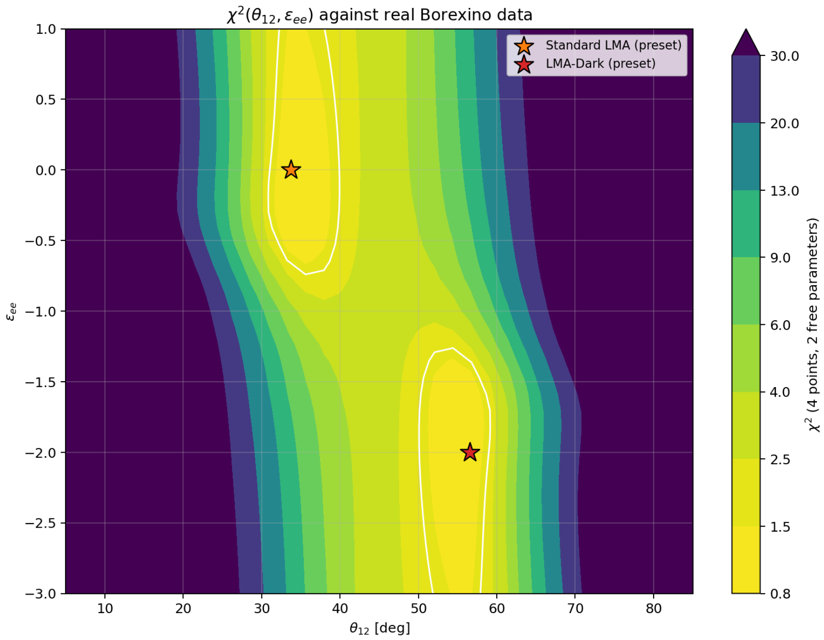

Figure 4.4: `inference2_borexino_nsi.ipynb`, `inference2_fig1_chi2_landscape.png`: the 30×30 $\Delta(-2\ln\mathcal{L})$ grid scan over $(\theta_{12}, \epsilon_{ee})$, with the standard-LMA and LMA-Dark preset points marked, the source of the confirmed grid-scan minimum above.

Approximations and Limitations

Confirmed, directly printed output. The exact numerical values of the LBFGS fits themselves (Section 4 of the notebook) do have recorded cell output in the saved file:

```
LMA-side fit (started near theta12=25 deg, eps_ee=+0.5):
fit    :  theta12=34.50 +/- 4.08 deg, eps_ee=-0.225 +/- 0.685
chi2   :  10.714 -> 0.737 (17 closure evals)
Dark-side fit (started near theta12=65 deg, eps_ee=-1.2):
fit    :  theta12=55.50 +/- 4.08 deg, eps_ee=-1.782 +/- 0.700
chi2   :  18.762 -> 0.737 (14 closure evals)
|chi2_LMA - chi2_dark| = 0.0002 - indistinguishable given the data
PASSED: both starting points converge to statistically equivalent
solutions.
```

This confirms, rather than merely back-derives, the Section 5 cross-check above: $\theta_{12}^{\mathrm{LMA}} = 34.50°$ and $\epsilon_{ee}^{\mathrm{LMA}} = -0.225$ match exactly the values that table's "Predicted from LMA fit" column was computed from, and the LMA/dark-side fit's own `assert delta_chi2 < 1.0` passes with $|\chi^2_{\mathrm{LMA}} - \chi^2_{\mathrm{dark}}| = 0.0002$. Both fits are also reasonably consistent with the independently confirmed grid-scan minimum, $\theta_{12} = 35.34°$, $\epsilon_{ee} = -0.24$ (a coarse 30×30 grid evaluated independently of the gradient fit), the expected agreement between two different optimisation methods.

The notebook produces three figures: a filled $\chi^2(\theta_{12}, \epsilon_{ee})$ contour map over the full grid with the standard-LMA and LMA-Dark preset points marked as reference stars; an overlay of the two best-fit $P_{ee}(E)$ curves against the real Borexino data points with a residual-difference panel; and the $\Delta(-2\ln\mathcal{L})$ uncertainty map with Wilks' 68%/90%/95% contours (at $\Delta\chi^2 = 2.30/4.61/6.18$ for 2 degrees of freedom) and the grid minimum marked. The notebook's own stated conclusion: "two deliberately different, uninformative starting guesses converge to two well-separated points...with statistically indistinguishable χ^2...quantitatively consistent with the $\theta_{12} \to 90° - \theta_{12}$, $\epsilon_{ee} \to -(2 + \epsilon_{ee})$ relation," and "the two fitted $P_{ee}(E)$ curves agree to a fraction of a percent everywhere from 0.15 to 16 MeV" – but also, echoing Section 4.2, that "a real resolution of LMA vs. LMA-Dark needs independent information...not more solar $P_{ee}(E)$ points alone" (i.e. the Earth day–night asymmetry of Section 4.2, or another observable outside the solar $P_{ee}(E)$ degeneracy itself).

4.5 The composition-dependence extension: no notebook coverage yet

Approximations and Limitations

This is the most important caveat in this chapter. An exhaustive search of every notebook in the repository (grep for `eps_ee_n`, `eps_.*_n`, `epsilon_n`, `has_neutron_coupling`, and `composition`) confirms that **no notebook currently demonstrates the neutron-density composition-dependence extension** (ϵ^n, Sections 1.3 and 2.3). The handful of unrelated matches (an Earth/solar/AK135 raw-data column literally named `neutron_density_mol_cm3`, evolutor-composition consistency checks unrelated to NSI, PMNS flavour-composition language, etc.) are all false positives on the search terms, not partial coverage. This is consistent with the extension's implementation history: it was added after every existing notebook was last executed, so no notebook artifact can yet be cited as numerical evidence for it. The automated test suite (`core/BSM/test/test2_bsm_nsi.py`, `core/common/test/test3_hamiltonian.py`, `core/perturbative/test/test2_perturbative_evolutor.py`, and the `medium/earth`, `medium/atmosphere`, `medium/solar` test suites) is, at the time of writing, the only executed and verified evidence for the composition-dependence extension's correctness; a natural follow-up to this document would be a fifth notebook demonstrating it end to end – e.g. propagating a solar neutrino from the hydrogen-rich envelope ($n_n/n_e \to 0$) through the helium/metal core ($n_n/n_e \to 1$) with a non-zero ϵ^n, and comparing against the $\epsilon^n = 0$ approximation used throughout this chapter's existing notebooks.

References

[1] S. K. Agarwalla, S. S. Chatterjee, and A. Palazzo. "Degeneracy between theta23 octant and neutrino non-standard interactions at DUNE". In: *Phys. Lett. B* 762 (2016), pp. 64–71. eprint: 1607.01745.

[2] C. Biggio, M. Blennow, and E. Fernandez-Martinez. "General bounds on non-standard neutrino interactions". In: *JHEP* 08 (2009), p. 090. eprint: 0907.0097.

[3] P. Coloma, P. B. Denton, M. C. Gonzalez-Garcia, M. Maltoni, and T. Schwetz. "Curtailing the Dark Side in Non-Standard Neutrino Interactions". In: *JHEP* 04 (2017), p. 116. eprint: 1701.04828.

[4] A. M. Dziewonski and D. L. Anderson. "Preliminary reference Earth model". In: *Phys. Earth Planet. Inter.* 25 (1981), pp. 297–356.

[5] I. Esteban, M. C. Gonzalez-Garcia, M. Maltoni, I. Martinez-Soler, and T. Schwetz. "Updated constraints on non-standard interactions from global analysis of oscillation data". In: *JHEP* 08 (2018), p. 180. eprint: 1805.04530.

[6] Y. Farzan and M. Tórtola. "Neutrino oscillations and Non-Standard Interactions". In: *Front. in Phys.* 6 (2018), p. 10. eprint: 1710.09360.

[7] S. Fukasawa and O. Yasuda. "The possibility to observe the non-standard interaction by the Hyperkamiokande atmospheric neutrino experiment". In: *Nucl. Phys. B* 914 (2017), pp. 99–116. eprint: 1608.05897.

[8] Y. Grossman. "Nonstandard neutrino interactions and neutrino oscillation experiments". In: *Phys. Lett. B* 359 (1995), pp. 141–147. eprint: hep-ph/9507344.

[9] IceCube Collaboration. "Search for Nonstandard Neutrino Interactions with IceCube DeepCore". In: *Phys. Rev. D* 97 (2018), p. 072009. eprint: 1709.07079.

[10] IceCube Collaboration. "All-flavor constraints on nonstandard neutrino interactions and generalized matter potential with three years of IceCube DeepCore data". In: *Phys. Rev. D* 106 (2022), p. 032009. eprint: 2106.07755.

[11] K. J. Kelly. "Searches for new physics at the Hyper-Kamiokande experiment". In: *Phys. Rev. D* 95 (2017), p. 115009. eprint: 1703.00448.

[12] T. Kikuchi, H. Minakata, and S. Uchinami. "Perturbation Theory of Neutrino Oscillation with Nonstandard Neutrino Interactions". In: *JHEP* 03 (2009), p. 114. eprint: 0809.3312.

[13] S. P. Mikheyev and A. Y. Smirnov. "Resonance enhancement of oscillations in matter and solar neutrino spectroscopy". In: *Sov. J. Nucl. Phys.* 42 (1985), pp. 913–917.

[14] T. Ohlsson. "Status of non-standard neutrino interactions". In: *Rept. Prog. Phys.* 76 (2013), p. 044201. eprint: 1209.2710.

[15] A. Renshaw et al. "First Indication of Terrestrial Matter Effects on Solar Neutrino Oscillation". In: *Physical Review Letters* 112 (2014), p. 091805. DOI: 10.1103/PhysRevLett.112.091805. eprint: 1312.5176.

[16] N. Vinyoles et al. "A New Generation of Standard Solar Models". In: *Astrophys. J.* 835 (2017), p. 202. eprint: 1611.09867.

[17] L. Wolfenstein. "Neutrino oscillations in matter". In: *Phys. Rev. D* 17 (1978), pp. 2369–2374.