

The 3+1 Sterile Neutrino Extension in `tpeanuts`

Fast Simulation of Neutrino Oscillations in Matter

MSc Thesis (Trabajo Fin de Máster) — BSM Extension Companion Document

Juan Ramon Diaz Santos
`diazjuan@alumni.uv.es`

Supervisors: Roberto Ruiz de Austri Bazan & Michele Lucente

August 2026

This document is the companion reference for the `core.BSM.bsm_sterile` and `core.BSM.bsm_mass_spectrum` packages and every downstream consumer of a 4-flavour `OscillationParameters.pmns`: the exact-Hamiltonian builders in `core.common.hamiltonian`, the closed-form Ferrari eigenvalue solver in `core.perturbative.spectral`, the numerical and perturbative evolutors, and the solar/Earth/atmosphere propagation pipelines. Chapter 1 covers the theoretical foundations of the 3+1 scheme: the short-baseline anomalies motivating it, the 4×4 mixing parametrisation, the sterile neutral-current term, and the exact SM-limit embedding. Chapter 2 is an exhaustive account of the `tpeanuts` implementation, including the closed-form quartic (Ferrari) eigenvalue solution specific to $N = 4$. Chapter 3 documents the named sterile presets and the observables/simulations that can be built with them. Chapter 4 walks through the validation notebooks that exercise the sterile extension and reports their actual results. Blue boxes (**Theoretical Foundation**) contain the physics; green boxes (**Implementation**) link that physics to actual `tpeanuts` code; orange boxes collect approximations and limitations; indigo boxes (**Result**) report a concrete numerical outcome of a notebook or of this project's own test suite. Text highlighted in `yellow` marks key results/equations.

Contents

1	Theoretical Foundations of the 3+1 Sterile Extension	1
1.1	Motivation: short-baseline anomalies	1
1.2	The 3+1 mixing matrix	2
1.3	The matter potential: charged current versus neutral current	2
1.4	The two-flavour disappearance limit	3
2	Implementation in tpeanuts	4
2.1	PMNSSterileParams and PMNS_sterile	4
2.2	MassSpectrum_BSM: the sterile mass splitting	5
2.3	Hamiltonian assembly for $N = 4$	5
2.4	The closed-form quartic (Ferrari) eigenvalue solver	6
2.5	The perturbative segment evolver	8
2.6	Pipeline wiring: the neutral-current policy	8
2.7	Summary: modules touched by the sterile extension	9
3	Presets, Observables, and Simulations	10
3.1	Named sterile presets	10
3.2	Building an oscillation scenario with the sterile extension	11
3.3	Observables and simulations that can be built with the sterile extension	11
4	Notebook Results	13
4.1	sterile1_test.ipynb: framework verification	13
4.2	sterile2_kinematics.ipynb: the two-flavour approximation	14
4.3	sterile3_analysis.ipynb: experimental anomalies and the global tension	15
4.4	intrinsic3_BSM_extension_sterile.ipynb: cross-method validation	17

Theoretical Foundations of the 3+1 Sterile Extension

This chapter develops the physics of the 3+1 sterile-neutrino extension, from the short-baseline anomalies that originally motivated it to the exact matter-potential structure a fourth, gauge-singlet mass eigenstate implies. It is deliberately implementation-agnostic: [Chapter 2](#) maps every equation derived here onto the concrete `tpeanuts` code that evaluates it.

1.1 Motivation: short-baseline anomalies

Theoretical Foundation

The three-flavour paradigm, with two independent mass-squared splittings $\Delta m_{21}^2 \sim 7.4 \times 10^{-5} \text{ eV}^2$ and $\Delta m_{31}^2 \sim 2.5 \times 10^{-3} \text{ eV}^2$, accounts for the vast majority of solar, atmospheric, reactor, and accelerator oscillation data. A handful of short-baseline anomalies, however, are not accommodated by this picture and instead point to oscillations at a much larger mass-squared splitting, $\Delta m^2 \sim \mathcal{O}(1) \text{ eV}^2$, requiring a third, independent splitting and hence (at least) a fourth mass eigenstate:

- **LSND:** excess $\bar{\nu}_\mu \rightarrow \bar{\nu}_e$ appearance at $L/E \sim 1 \text{ m/MeV}$ [1].
- **MiniBooNE:** an independent, statistically significant excess of electron-like events in both neutrino and antineutrino mode, broadly consistent with the LSND signal [2].
- **The reactor antineutrino anomaly:** a $\sim 6\%$ deficit between measured and predicted $\bar{\nu}_e$ flux at short baseline from reactor experiments, once reactor flux predictions were revised [8].
- **The gallium anomaly:** a deficit in the ν_e event rate measured during the radioactive-source calibration campaigns of the GALLEX and SAGE solar-neutrino experiments [5].

The simplest new-physics explanation common to all four is a single additional mass eigenstate ν_4 , mostly composed of a new flavour eigenstate ν_s that does not participate in the weak interaction at all (a “sterile” neutrino) – the “3+1” scheme. The appearance anomalies (LSND, MiniBooNE) and the disappearance anomalies (reactor, gallium) probe different combinations of the new mixing angles, and existing global fits are in tension between the two classes; this tension, and the different individual mixing angles preferred by different data sets, is exactly why [Chapter 3](#) offers several distinct named presets rather than a single best fit.

1.2 The 3+1 mixing matrix

Theoretical Foundation

A fourth mass eigenstate ν_4 and a fourth sterile flavour eigenstate ν_s are added to the three active ones. Because ν_s is a gauge-singlet, it does not receive the charged-current matter potential of §1.3 below, and the 4×4 mixing matrix is parametrised, following the convention of Kopp et al. [7], by adding three active–sterile rotations R_{14}, R_{24}, R_{34} to the standard 3-flavour structure,

$$U_4 = R_{23} \Delta_4 R_{24} R_{34} R_{13} \Delta_4^\dagger R_{12} R_{14}, \quad (1.1)$$

where $\theta_{14}, \theta_{24}, \theta_{34}$ are the new active–sterile mixing angles and δ_{14}, δ_{24} their associated CP phases. Equation (1.1) reduces exactly to the 3-flavour Standard Model matrix (*embedded* in the larger 4×4 space) when $\theta_{14} = \theta_{24} = \theta_{34} = 0$: the sterile row/column decouple completely and every propagation result of the pure Standard Model is recovered without approximation.

Approximations and Limitations

CP-phase counting. For N Dirac generations there are $(N-1)(N-2)/2$ independent physical CP phases. For $N = 4$ this gives exactly **three**: the standard phase δ_{13} (carried through R_{13} by Δ_4 , exactly as in the 3-flavour case), the e – s phase δ_{14} (carried by R_{14}), and the μ – s phase δ_{24} (carried by R_{24}). R_{34} is always a *real* rotation: the fourth possible phase can be absorbed into a redefinition of the charged-lepton fields, so no sterile preset (Chapter 3) ever needs a δ_{34} value.

1.3 The matter potential: charged current versus neutral current

Theoretical Foundation

Because ν_s is a gauge singlet, it never receives the charged-current potential: in the extended flavour basis (e, μ, τ, s) , the CC term is simply the 3-flavour Wolfenstein potential [10], zero-padded,

$$H_{\text{mat}}^{\text{CC}} = \text{diag}(V_{CC}, 0, 0, 0), \quad V_{CC} = \pm\sqrt{2} G_F n_e L_{\text{scale}}. \quad (1.2)$$

The three *active* flavours, in contrast, all share the same neutral-current potential $V_{NC} = \mp \frac{1}{\sqrt{2}} G_F n_n L_{\text{scale}}$ equally (lepton universality of the NC coupling). In the pure 3-flavour Standard Model this shared term is an unobservable global phase and is correctly never computed at all. Once a sterile state is present – an NC-blind fourth entry – it is no longer a pure phase: writing out the full diagonal and subtracting the common piece,

$$\text{diag}(V_{CC} + V_{NC}, V_{NC}, V_{NC}, 0) - V_{NC} I_4 = \text{diag}(V_{CC}, 0, 0, -V_{NC}), \quad (1.3)$$

a genuine, physically observable relative potential of magnitude comparable to V_{CC} itself survives on the sterile diagonal entry. This term is *optional* in the sense that omitting it (the default) recovers the CC-only sterile Hamiltonian used throughout most of the sterile-oscillation literature and throughout earlier phases of this project – unlike the purely 3-flavour case, though, omitting it is a genuine approximation, not an exact

simplification, once θ_{24} or θ_{34} (which mix the sterile index into the μ/τ sector via the outer block, [Chapter 2](#)) are non-zero.

Approximations and Limitations

This sterile neutral-current term is structurally independent of, but can be combined freely with, Non-Standard Interactions (a separate BSM extension; see the companion `BSM_NSI` document): both consume the same underlying local neutron density $n_n(x)$, but they are additive, independently switchable contributions to the matter Hamiltonian.

1.4 The two-flavour disappearance limit

Theoretical Foundation

When every Standard Model active mixing angle is switched off ($\theta_{12} = \theta_{13} = \theta_{23} = 0$), the ν_e survival probability reduces exactly to the textbook two-flavour vacuum formula

$$P_{ee} = 1 - \sin^2(2\theta_{14}) \sin^2\left(1.267 \frac{\Delta m_{41}^2 [\text{eV}^2] L [\text{km}]}{E [\text{GeV}]}\right), \quad (1.4)$$

for *any* value of θ_{24} , since θ_{24} lives entirely in the outer block O_4 ([Chapter 2](#)), which leaves the electron index untouched. The analogous statement for ν_μ disappearance, $|U_{\mu 4}|^2 = \sin^2 \theta_{24}$, holds exactly only when *both* $\theta_{23} = 0$ and $\theta_{14} = 0$: R_{14} and R_{24} both act on the sterile index and do not commute in general, so a small residual correction survives in the ν_μ channel whenever θ_{14} and θ_{24} are simultaneously non-zero – a property of the physics, not of the numerical implementation, and the subject of [Section 4.2](#)'s quantitative study.

Key References

Short-baseline anomalies: [\[1, 2, 5, 8\]](#). 3+1 mixing parametrisation: [\[7\]](#). Global 3+1 fit: [\[4\]](#). IceCube atmospheric ν_μ disappearance search: [\[6\]](#). Standard matter potential: [\[10\]](#).

Chapter 2

Implementation in tpeanuts

This chapter maps every equation of [Chapter 1](#) onto the concrete `tpeanuts` implementation, module by module. As with NSI (see the companion `BSM_NSI` document), there is no separate $N = 4$ -specific Hamiltonian-assembly module: `hamiltonian_reduced` and `hamiltonian_flavour` in `core.common.hamiltonian` handle the 3+1 extension by reading the flavour count directly off `OscillationParameters.pmns`, with no explicit branch on “is this sterile.”

2.1 PMNSSterileParams and PMNS_sterile

`core/BSM/bsm_sterile.py`

PMNS_sterile: the 4×4 mixing matrix for one additional light sterile neutrino.

Implementation in tpeanuts

`PMNSSterileParams` is an immutable dataclass holding *only* the sterile-sector extension: $\theta_{14}, \theta_{24}, \theta_{34}$ and their CP phases $\delta_{14}, \delta_{24}, \delta_{34}$ (the last always 0, [Chapter 1](#)). The Standard Model sector ($\theta_{12}, \theta_{13}, \theta_{23}, \delta_{13}$) lives in a separate, already-existing `PMNSParams` object, passed alongside it to the `PMNS_sterile` constructor – mirroring exactly how `NSIConfig` is an optional, orthogonal attribute of `OscillationParameters` rather than baked into the mixing object. Δm_{41}^2 is deliberately *not* stored on `PMNSSterileParams`: like $\Delta m_{21}^2 / \Delta m_{3\ell}^2$, it never enters a mixing-matrix rotation, so it lives on `OscillationParameters` (via `MassSpectrum_BSM`, below) instead.

`PMNS_sterile` is a direct subclass of the abstract `core.common.pmns.PMNS` with the class attribute `n_flavours=4`: every method that is generic in flavour count (`R12`, `R13`, `R23`, `Delta`, `flavour_basis`, `select_antinu`, ...) is inherited *unchanged*, automatically sized to 4×4 ; only the genuinely 3+1-specific pieces are defined here.

Theoretical Foundation

The three new rotation builders each act on a single active–sterile 2-plane, mirroring `R12/R13/R23`’s construction exactly (via the shared `PMNS._rot` generator):

`R14()` = rotation in the $(e, s) = (0, 3)$ plane, carrying δ_{14} ,

`R24()` = rotation in the $(\mu, s) = (1, 3)$ plane, carrying δ_{24} ,

`R34()` = rotation in the $(\tau, s) = (2, 3)$ plane, real ($\delta_{34} = 0$).

`outer_block` and `reduced` are overridden with the 3+1 product structure of [Eq. \(1.1\)](#), identifying $O_4 \equiv R_{23}\Delta_4 R_{24}R_{34}$ (`PMNS_sterile.outer_block`) so that a formal reduced factor $U_{\text{red},4} \equiv R_{13}\Delta_4^\dagger R_{12}R_{14}$ satisfies $U_4 = O_4 U_{\text{red},4}$ exactly, mirroring the 3-flavour Standard Model factorisation. Exactly as for Δ in the 3-flavour case, Δ_4 cancels from the reduced *kinetic term* specifically,

$$\Delta_4^\dagger (R_{12}R_{14} \text{diag}(k) R_{14}^\dagger R_{12}^\dagger) \Delta_4 = R_{12}R_{14} \text{diag}(k) R_{14}^\dagger R_{12}^\dagger, \quad (2.1)$$

valid because neither R_{12} (the e – μ plane) nor R_{14} (the e – s plane) mixes with the τ index, the only non-trivial index of Δ_4 . This is why `PMNS_sterile.reduced` returns the simplified, Δ_4 -free

$$U_{\text{red},4}^{\text{code}} = R_{13}R_{12}R_{14} \quad (2.2)$$

rather than the formal $U_{\text{red},4}$: the two agree once contracted with the diagonal kinetic term, but $U_{\text{red},4}^{\text{code}}$ is cheaper to build and use.

Approximations and Limitations

$U = O U_{\text{red}}$ holds exactly only for the formal U_{red} , not for `reduced`. Consequently, `outer_block().reduced()` is *not* bit-for-bit equal to `pmns_matrix()`: the two agree up to a diagonal phase (Δ_4) that leaves the kinetic term invariant but is otherwise dropped. This is harmless for every use of `reduced` in this project (always as the sandwiching matrix of a kinetic term), and mirrors the identical, already-documented situation for the plain 3-flavour `PMNS_SM`.

2.2 MassSpectrum_BSM: the sterile mass splitting

`core/BSM/bsm_mass_spectrum.py`

Extends the active-sector mass-squared-difference vector with Δm_{41}^2 .

Implementation in tpeanuts

`MassSpectrum_BSM` is a frozen dataclass carrying one extra field, `DeltamSq41`, on top of the inherited `DeltamSq21/DeltamSq31`. Its `difference_vector()` override appends `DeltamSq41` to the base 3-vector, returning a $(\dots, 4)$ -shaped tensor consumed by `kinetic_eigenvalue_vector` (Section 2.3) to build the four dimensionless kinetic eigenvalues $k_1, \dots, k_4$. `DeltamSq41=None` is accepted at construction time (so an incomplete configuration can still be built and inspected) but raises `ValueError` the moment `difference_vector` is actually evaluated, since a 3+1 Hamiltonian genuinely cannot be assembled without it.

2.3 Hamiltonian assembly for $N = 4$

`core/common/hamiltonian.py`

`hamiltonian_kinetic_reduced/hamiltonian_matter_reduced`: the same two functions used for the 3-flavour and NSI cases, generalised to any flavour count.

Theoretical Foundation

The reduced-basis kinetic term uses $U_{\text{red},4}^{\text{code}} = R_{13}R_{12}R_{14}$ together with its *Hermitian conjugate* (never the plain transpose):

$$H_{\text{kin},4}^{\text{red}} = U_{\text{red},4}^{\text{code}} \text{diag}(k_1, k_2, k_3, k_4) (U_{\text{red},4}^{\text{code}})^\dagger. \quad (2.3)$$

Unlike the 3-flavour $R_{13}R_{12}$, $U_{\text{red},4}^{\text{code}}$ is genuinely complex whenever the active–sterile

CP phase $\delta_{14} \neq 0$ (it carries R_{14} , built with a non-trivial phase, unlike R_{12}/R_{13}): the Hermitian-conjugate convention is therefore not a merely-safe default here but the only correct choice for $U_4 = O_4 U_{\text{red},4}$ to be respected by the flavour-basis dressing.

Theoretical Foundation

The matter term takes one of three forms, matching [Section 1.3](#)’s three cases exactly:

1. **No NSI, no neutral-current term (the default).** Since $R_{23}, \Delta_4, R_{24}, R_{34}$ (i.e. every generator of O_4) leave the electron index fixed, $O_4 \text{diag}(V_{CC}, 0, 0, 0) O_4^\dagger = \text{diag}(V_{CC}, 0, 0, 0)$ exactly, so no rotation is needed at all: $H_{\text{mat},4}^{\text{red}} = H_{\text{mat},4}^{\text{flavour}} = \text{diag}(V_{CC}, 0, 0, 0)$. This introduces no approximation beyond the 3-flavour case, for any $\theta_{14}, \theta_{24}, \theta_{34}$.
2. **With the neutral-current term, no NSI.** Supplying a neutron density activates $H_{\text{mat},4}^{\text{flavour}} = \text{diag}(V_{CC}, 0, 0, -V_{NC})$ ([Eq. \(1.3\)](#)) instead. Because O_4 genuinely mixes the $\{\mu, \tau, s\}$ block, it no longer leaves this term invariant, so `hamiltonian_matter_reduced` takes the general rotated path,

$$H_{\text{mat},4}^{\text{red}} = O_4^\dagger H_{\text{mat},4}^{\text{flavour}} O_4, \quad (2.4)$$

even though NSI is off – the reduction-trick shortcut of case 1 is specific to the charged-current-only term.

3. **With NSI (and, optionally, the neutral-current term).** Fully supported by the same code path documented in the companion BSM_NSI document: `NSIConfig.epsilon_tensor(n_flavours=4)` (and `epsilon_n_tensor`) embed the 3×3 active-flavour block in the top-left corner of a 4×4 zero matrix, since ν_s is a gauge singlet and carries no NSI coupling of either kind; the resulting flavour-basis matter term is rotated into the reduced basis with the same $O_4^\dagger(\cdot)O_4$ transformation as case 2.

Nothing in `hamiltonian_reduced/hamiltonian_flavour` branches explicitly on “sterile or NSI or both”: the flavour count $n \in \{3, 4\}$ and the presence of `oscillation.nsi` are read directly off `OscillationParameters`, so all combinations are handled by the exact same two functions used throughout this project.

2.4 The closed-form quartic (Ferrari) eigenvalue solver

`core/perturbative/spectral.py`

`hamiltonian_traceless_eigenvalues(..., analytic=True)`: closed-form eigenvalues for a traceless Hermitian Hamiltonian, Cardano for $N = 3$, Ferrari for $N = 4$.

Theoretical Foundation

The perturbative evolver ([Section 2.5](#)) needs the Hamiltonian’s spectral decomposition $H = \sum_a \lambda_a M_a$. By default λ_a is obtained numerically via `torch.linalg.eigvalsh`; an explicit closed-form alternative is also available specifically for $N \in \{3, 4\}$, since a traceless Hermitian matrix’s characteristic polynomial is guaranteed real-rooted and has no sub-leading term. For $N = 4$ specifically – the case this document is about – the

characteristic polynomial of the traceless part $T = H - \text{Tr}(H)I_4/4$ is the depressed quartic $\lambda^4 + p\lambda^2 + q\lambda + r = 0$ with $p = c_1 = -\frac{1}{2}\text{Tr}(T^2)$, $q = -e_3 = -\frac{1}{3}\text{Tr}(T^3)$, $r = e_4 = \det(T)$. Ferrari's method solves it by first finding a real root m of the *resolvent cubic*

$$m^3 + pm^2 + \left(\frac{p^2}{4} - r\right)m - \frac{q^2}{8} = 0, \quad (2.5)$$

itself solved with the same trigonometric (Cardano) formula used for the $N = 3$ case: since T is Hermitian, the quartic is guaranteed real-rooted, and the resolvent cubic's three roots are exactly the real pairwise sums $m_{ij|kl} = -\frac{1}{2}(\lambda_i + \lambda_j)(\lambda_k + \lambda_l)$ over the three ways of splitting the four roots into two pairs, so it is real-rooted too and the closed form applies unconditionally. Taking the *largest* resolvent root m (the conditioning-optimal choice) factors the quartic into two real quadratics,

$$s = \sqrt{2m}, \quad \lambda \in \left\{ \frac{s \pm \sqrt{\Delta_1}}{2}, \frac{-s \pm \sqrt{\Delta_2}}{2} \right\}, \quad \Delta_1 = -2\left(m + p + \frac{q}{s}\right), \quad \Delta_2 = -2\left(m + p - \frac{q}{s}\right). \quad (2.6)$$

Approximations and Limitations

Conditioning is not unconditional for $N = 4$, unlike $N = 3$. The chosen pairing's m can itself be close to singular (only possible, for a traceless Hermitian T , near $T \approx 0$), or the resolvent cubic's own three roots can nearly coincide even when the *chosen* m is well-separated from zero (e.g. $\lambda = (a, a, -a, -a)$ has a well-conditioned pairing $m = 2a^2$ whose resolvent cubic nonetheless has a repeated root elsewhere). Batch entries flagged by either check fall back directly to `torch.linalg.eigvalsh` for the *eigenvalues* themselves (unlike the projector-only degeneracy fallback described next), since this closed form can lose the eigenvalues, not merely lose precision in a downstream ratio.

Spectral projectors and the degeneracy fallback. The corresponding $N = 4$ projector formula,

$$M_a = \frac{T^3 + \lambda_a T^2 + (c_1 + \lambda_a^2)T + (\lambda_a^3 + c_1 \lambda_a - e_3)I_4}{4\lambda_a^3 + 2c_1 \lambda_a - e_3}, \quad (2.7)$$

has a denominator equal to $p'(\lambda_a)$, the characteristic polynomial's derivative at that root, which vanishes exactly at a repeated eigenvalue. Near a (near-)degenerate spectrum, `_spectral_degeneracy_mask` flags the affected batch entries (by relative eigenvalue-gap size) and replaces *only their projectors* wholesale with $M_a = v_a v_a^\dagger$ from a direct `torch.linalg.eigh(T)` call – exact and stable by construction, computed only when at least one batch entry actually needs it.

Implementation in tpeanuts

`hamiltonian_traceless_eigenvalues(T, *, analytic=False, c1=None, c0=None, e3=None, e4=None)` dispatches on `T.shape[-1]`: Cardano for 3, Ferrari for 4, and unconditional fallback to `torch.linalg.eigvalsh` for any other size (no closed form is defined there). `evolutor_perturbative_segment` and `earth_evolver/atmosphere_evolver` expose this as the public `analytic_eigenvalues=True` flag; passing it together with a matrix-exponential (`method="numerical"`) propagation path raises, since that path never diagonalises the

Hamiltonian at all.

2.5 The perturbative segment evolver

`core/perturbative/evolver.py`

evolver_perturbative_segment: generalised, without modification, to $N = 4$.

Implementation in tpeanuts

`evolver_perturbative_segment` never branches explicitly on “sterile”: it reads `n_flavours = int(pmns.n_flavours)` off whatever **PMNS** object **OscillationParameters** carries and sizes every subsequent tensor accordingly. The one behaviour that *is* sterile-specific is `include_matter_nc`: when `True` (and `n_flavours==4`), the sterile neutral-current direction matrix P_{nc} (Eq. (1.3), rotated by $O_4^\dagger(\cdot)O_4$ exactly as in Section 2.3’s case 2) is built by `_matter_direction_reduced` and sandwiched into the first-order correction alongside the CC direction P , using the profile’s `residual_integral_neutron(1a, 1b)` and `matter_potential_nc`. Requesting it without a profile built with neutron-density coefficients (`potential_n/average_n/ residual_integral_neutron` unavailable) raises `ValueError` rather than silently omitting the term.

2.6 Pipeline wiring: the neutral-current policy

Approximations and Limitations

resolve_include_matter_nc: the shared tri-state policy. Because omitting the sterile NC term is a genuine physical incompleteness once θ_{24} or θ_{34} are non-zero (Section 1.3), `medium.solar`, `medium.earth`, and `medium.atmosphere` all resolve `include_matter_nc` through one shared function in `core.common.oscillation`:

- `include_matter_nc` given explicitly (`True/False`): always honoured as-is.
- `include_matter_nc=None` (unspecified): resolves to `True` when `oscillation.BSM_extension_sterile` is set *and* the profile/medium in hand can actually supply a neutron density; `False` for the plain 3-flavour case regardless (there V_{NC} truly is an unobservable common phase, so omitting it is exact, not approximate). When sterile is active but no neutron-density data is available, this still resolves to `False` (the CC-only fallback), but emits an explicit `RuntimeWarning` rather than resolving silently – exactly the warning quoted in Section 4.4, triggered because the `even_power` Earth profile used by that notebook’s Section 5 was not built with `include_neutron=True`.

Each medium then supplies the raw density itself: `EarthProfile` via `profile_perturbative_kwargs='include_neutron': True` at construction time (PREM- or even-power-tabulated), `AtmosphereParameters` via its own `include_matter_nc` field feeding `atmosphere_density(..., density_type="neutron_density")`, and `SolarMediumProfile` via its `density_n` table (derived from the standard solar model composition).

2.7 Summary: modules touched by the sterile extension

Module	Role
<code>core/BSM/bsm_sterile.py</code>	PMNS_sterile : 4×4 mixing matrix, $R_{14}/R_{24}/R_{34}$
<code>core/BSM/bsm_mass_spectrum.py</code>	MassSpectrum_BSM : appends Δm_{41}^2
<code>core/common/hamiltonian.py</code>	hamiltonian_reduced/_flavour , generic in flavour count
<code>core/perturbative/spectral.py</code>	Ferrari closed-form quartic eigenvalues, $N = 4$
<code>core/perturbative/evolutor.py</code>	<code>P_nc</code> sandwich term, analytic_eigenvalues
<code>core/numerical/evolutor.py</code>	Exact matrix-exponential evolution, generic in flavour count
<code>core/common/oscillation.py</code>	resolve_include_matter_nc : shared tri-state policy
<code>medium/earth</code> , <code>medium/atmosphere</code> , <code>medium/solar</code>	Neutron-density profile construction/sampling

Presets, Observables, and Simulations

3.1 Named sterile presets

config/presets.py

OSCILLATION_PRESETS entries carrying a `theta14_deg` key: named, literature-justified 3+1 parameter sets, dispatched to **PMNS_sterile** automatically.

Implementation in tpeanuts

A preset is a *sterile* preset purely by having a `theta14_deg` key; `PropagationConfig.oscillation_parameters_from_preset` dispatches on that key alone to build a **PMNS_sterile** instead of a **PMNS_SM**, with every other field (the SM angles, $\Delta m_{21}^2/\Delta m_{3\ell}^2$) inherited from the shared NuFIT 5.2 baseline [3, 9]. As with the NSI presets of the companion document, these are named, traceable benchmark points rather than a single global best fit: current 3+1 data is in tension between different anomalies (Chapter 1), so no universally preferred point exists.

Preset	Parameters and origin
<code>sterile_3p1_null_mixing</code>	$\theta_{14} = \theta_{24} = \theta_{34} = 0$, $\Delta m_{41}^2 = 1.0 \text{ eV}^2$. Null-hypothesis reference: with every sterile angle at zero, PMNS_sterile is numerically identical to the plain 3-flavour SM run through the 4-flavour machinery – validated to $\sim 10^{-16}$ in Section 4.1.
<code>sterile_3p1_bestfit_giunti2017</code>	$\theta_{14} = 8.5^\circ$, $\theta_{24} = 7.5^\circ$, $\theta_{34} = 0$, $\Delta m_{41}^2 = 1.7 \text{ eV}^2$ ($\sin^2 2\theta_{14} = 0.085$, $\sin^2 2\theta_{24} = 0.068$). Global 3+1 fit combining the reactor and gallium anomalies with solar data [4]; sterile CP phases fixed to zero (no sensitivity in that data combination). The primary benchmark used throughout Chapter 4.
<code>sterile_3p1_benchmark_icecube</code>	$\theta_{14} = 0$, $\theta_{24} \approx 9.22^\circ$ ($\sin^2 2\theta_{24} = 0.10$), $\theta_{34} = 0$, $\Delta m_{41}^2 = 0.3 \text{ eV}^2$. IceCube atmospheric ν_μ disappearance benchmark [6]; θ_{14}, θ_{34} set to zero since IceCube atmospheric neutrinos are primarily sensitive to θ_{24} .
<code>sterile_3p1_two_flavour_limit</code>	$\theta_{12} = \theta_{13} = \theta_{23} = 0$, $\theta_{14} = 15^\circ$, $\theta_{24} = 12^\circ$, $\theta_{34} = 0$, $\Delta m_{41}^2 = 1.0 \text{ eV}^2$. A mathematical construction, not a data-driven benchmark: switching off every SM active angle makes the ν_e disappearance channel reduce exactly to the textbook 2-flavour formula for any θ_{24} . Used in Section 4.2 to isolate the 2-flavour approximation from the active-sector suppression quantified there for the realistic presets above.

Key References

[4, 6, 7].

3.2 Building an oscillation scenario with the sterile extension

Attaching a sterile preset to a propagation config

```
from tpeanuts.config.propagation import PropagationConfig
from tpeanuts.util.context import RuntimeContext

context = RuntimeContext.resolve("cpu", torch.float64)
oscillation = PropagationConfig.oscillation_parameters_from_preset(
    "sterile_3p1_bestfit_giunti2017",  # any theta14_deg-carrying preset
    context=context,
)
# oscillation.pmns is now a PMNS_sterile (n_flavours == 4), and
# oscillation.mass_spectrum carries DeltamSq41; every downstream builder
# (hamiltonian_reduced/flavour, evolutor_numerical,
# evolutor_perturbative_segment, medium.*) generalises automatically.
```

Enabling the sterile neutral-current term

```
from tpeanuts.medium.earth.probability import earth_probability_state

# include_matter_nc=None (default) auto-resolves to True here, since the
# preset is sterile and the profile below carries neutron-density data.
P = earth_probability_state(
    nustate, profile, oscillation, E_MeV, eta, depth_m,
    method="numerical", include_matter_nc=None,
)
```

3.3 Observables and simulations that can be built with the sterile extension

Implementation in tpeanuts

Every observable `tpeanuts` exposes for the Standard Model can be re-evaluated with a `PMNS_sterile` in place of `PMNS_SM`, subject to the method-compatibility notes below:

- **Vacuum ν_e/ν_μ disappearance.** The 4-flavour generalisation of the standard vacuum oscillation probability, used throughout [Chapter 4](#) to compare against the textbook 2-flavour formula and to locate short-baseline experiments in the (L, E) plane.
- **Solar $P_{ee}(E)$.** `medium.solar.probability.solar_probability_mass/_state` with `method="numerical"` or `method="adiabatic_exact"` (any N , including $N = 4$); the closed-form `method="adiabatic_approximated"` two-level approximation remains 3-flavour-only, exactly as for NSI.
- **Earth regeneration.** `medium.earth.probability.earth_probability_state/_numerical/_analytical`, with the sterile neutral-current term (§1.3) available via `include_matter_nc` whenever the Earth profile was built with neutron-density coefficients.
- **Atmospheric propagation.** `medium.atmosphere.evolutor.atmosphere_evolutor` and `atmosphere_probability_transition`, for ν_μ disappearance oscillograms directly comparable to the IceCube DeepCore benchmark.

- **Combined with NSI.** Any of the above with `oscillation.nsi` also set (companion `BSM_NSI` document) – the two BSM extensions compose freely, as documented in [Section 2.3](#)’s case 3.

Approximations and Limitations

As for NSI, `method="adiabatic_approximated"` (the closed-form two-level solar approximation and its Landau-Zener correction) rejects any BSM extension outright, including the sterile one, since its matter-mixing-angle formulas have no $N = 4$ generalisation. `analytic_eigenvalues=True` ([Section 2.4](#)) is available for both the perturbative Earth-/atmosphere paths at $N = 4$, but is incompatible with `method="numerical"`, which never diagonalises the Hamiltonian at all.

Notebook Results

This chapter reports the actual results produced by the four notebooks that exercise the 3+1 sterile extension: three physics-validation notebooks (`notebooks/validation/physics/BSM/sterile1_test.ipynb`, `sterile2_kinematics.ipynb`, `sterile3_analysis.ipynb`) and one cross-method intrinsic-validation notebook (`notebooks/validation/intrinsic/intrinsic3_BSM_extension_sterile.ipynb`). Unlike the companion BSM_NSI document, no inference notebook in this repository currently fits real data with the sterile extension.

Approximations and Limitations

On the completeness of the reported numbers. As with the companion BSM_NSI document's Chapter 4, several of these notebooks' saved `.ipynb` files have code cells with `execution_count: null` and an empty output list even though later cells or markdown narrative reference a result. This chapter states explicitly, for every quoted number, whether it is a value directly recorded in the notebook's saved output (**confirmed**) or not. This is a bookkeeping limitation of the notebook files as currently saved, not a defect in the underlying `tpeanuts` computation itself, which [Chapter 2](#) documents and this project's automated test suite (`core/BSM/test/test3_bsm_sterile.py` and related) independently exercises on every change.

4.1 `sterile1_test.ipynb`: framework verification

Implementation in `tpeanuts`

This notebook validates the 3+1 implementation against analytic expectations: the Standard-Model limit, unitarity, the matter-Hamiltonian structure, and CP-phase dependence, for three presets (`sterile_3p1_null_mixing`, `sterile_3p1_bestfit_giunti2017`, `sterile_3p1_benchmark_icecube`) plus the plain 3-flavour "`_SM_NUFIT52_NO`" reference, all propagated in vacuum via a local helper built from `hamiltonian_reduced`-style construction and `kinetic_potential` rather than through the Earth/atmosphere/solar pipelines.

Result

Confirmed, directly printed output.

SM-limit recovery ([Section 3.1](#)'s `sterile_3p1_null_mixing` vs. the plain 3-flavour SM): the active-sector column $|U_{\alpha 4}|^2$ for the null-mixing preset is exactly $[0, 0, 0, 1]$ (the sterile state fully decoupled); the PMNS submatrix agreement is $\|U_{\text{sterile}}[:, 3] - U_{\text{SM}}\|_F = \|\cdot\|_\infty = 0.00 \times 10^0$ (exact to the printed precision); the Hamiltonian agreement $\|H_4[:, 3] - H_{\text{SM}}\|_F$ is exactly 0 at all four tested (E, n_e) points (5 MeV/0.01, 10 MeV/0.1, 100 MeV/0.5, 1000 MeV/0.2 mol/cm³); the vacuum-probability agreement $\|P_4[:, 3] - P_3\|_\infty$ at $E = 10$ MeV is 2.22×10^{-16} (1 km), 3.47×10^{-18} (10 km), 2.00×10^{-15} (100 km), 0.00 (1000 km), 4.17×10^{-13} (10 000 km) – all consistent with floating-point round-off, not a

physical discrepancy. The notebook's own printed conclusion: "All diffs $\leq 10^{-6} \Rightarrow$ SM limit is correctly recovered."

Unitarity and probability conservation: $\|U_4 U_4^\dagger - I_4\|_F = 1.47 \times 10^{-16}$ (null mixing), 2.35×10^{-16} (Giunti 2017), 3.01×10^{-16} (IceCube 2020). Total vacuum probability at $E = 10$ MeV, $L = 100$ km: $\sum_\beta P_\beta = 1.00000000$ for all three presets, with maximum deviation $3.33 \times 10^{-16}/2.22 \times 10^{-16}/2.22 \times 10^{-16}$ respectively.

Sterile state receives no matter potential: for the Giunti-2017 preset at $E = 10$ MeV, $n_e = 0.1$ mol/cm³, both the sterile row and column of the matter Hamiltonian are exactly $[0+0i, 0+0i, 0+0i, 0+0i]$.

CP-phase scan ($\delta_{14} \in [0, 2\pi]$, custom `PMNSSterileParams` at LSND kinematics, $E = 10$ MeV, $L = 30$ m): "Max $|A_{CP}| = 0.001\%$ at $\delta_{14} = 108.9^\circ$," and disappearance CP-symmetry ($|P_{\mu\mu}(\nu) - P_{\mu\mu}(\bar{\nu})| \approx 0$) verified, consistent with the standard result that CP violation appears only in appearance channels.

Approximations and Limitations

The appearance CP asymmetry found here, $\lesssim 0.01\%$ at LSND kinematics, is tiny not because the framework lacks CP violation but because a short single baseline like LSND's cannot resolve two independent mass-squared splittings simultaneously – genuine 3+1 CP violation requires sensitivity to the interference between the standard and sterile oscillation frequencies, which needs either a longer baseline or an energy range spanning both scales. This is a property of the physics being probed, not a limitation of the numerical framework, and is exactly the kind of statement the notebook's own printed unitarity/SM-limit checks (both at floating-point precision) are designed to isolate from.

The notebook produces three figures: Fig 3, a three-panel heatmap of $|U_{\alpha i}|^2$ (flavour rows e, μ, τ, s against mass columns 1–4, annotated with numeric values) for the three presets; Fig 6, a three-panel heatmap of the reduced kinetic, matter, and total Hamiltonian for the Giunti-2017 preset; and Fig 9, $P(\nu_\mu \rightarrow \nu_e)/P(\bar{\nu}_\mu \rightarrow \bar{\nu}_e)/P(\nu_\mu \rightarrow \nu_\mu)/P(\nu_e \rightarrow \nu_e)$ and the CP asymmetry $A_{CP}(\%)$ against $\delta_{14} \in [0^\circ, 360^\circ]$.

4.2 `sterile2_kinematics.ipynb`: the two-flavour approximation

Implementation in `tpeanuts`

This notebook quantifies exactly the effect anticipated in Chapter 1's closing section: how well the textbook 2-flavour disappearance formula approximates the full 4-flavour vacuum evolution, separating the residual into (i) a pure mass-hierarchy ("one-mass-scale-dominance") piece and (ii) an active-sector mixing piece from R_{14}/R_{24} non-commutation, and mapping representative experiments onto the (L, E) plane.

Result

Confirmed, directly printed output – idealised limit (sterile_3p1_two_flavour_limit, $\theta_{12} = \theta_{13} = \theta_{23} = 0$, $\theta_{14} = 15^\circ$, $\theta_{24} = 12^\circ$, $\Delta m_{41}^2 = 1.0$ eV², compared against the exact-phase-constant 2-flavour formula):

Channel	max. $ P_{\text{numerical}} - P_{\text{2-flavour}} $
ν_e (both θ_{14}, θ_{24} active)	1.03×10^{-14}
ν_μ (θ_{14} forced to 0, isolated)	1.41×10^{-3}
ν_μ (both θ_{14}, θ_{24} active)	1.07×10^{-2}

Confirmed, directly printed output – realistic mixing (literature 1.267 phase constant):

Channel (preset)	$ U_{\alpha 4} ^2_{\text{naive}}$	$ U_{\alpha 4} ^2_{\text{exact}}$	predicted max. dev.	observed max. dev.
ν_e (Giunti 2017)	0.0218	0.0149	0.0268	0.0268
ν_μ (IceCube 2020)	0.0257	0.0110	0.0563	0.0821

Confirmed, directly printed output – oscillation-maximum kinematics: $\Delta m_{41}^2 = 0.3 \text{ eV}^2 \Rightarrow (L/E)_{\text{max}} = 4.13 \text{ m/MeV}$; $1.7 \text{ eV}^2 \Rightarrow 0.73 \text{ m/MeV}$; $7.0 \text{ eV}^2 \Rightarrow 0.18 \text{ m/MeV}$.

Approximations and Limitations

Reading the two confirmed tables together. The ν_e channel is exact to 10^{-14} (floating-point) in the idealised limit and its realistic-mixing deviation (0.0268) is *fully* explained by the naive-versus-exact $|U_{e4}|^2$ convention alone (predicted equals observed to the printed precision) – i.e. for ν_e , using $|U_{e4}|^2 = \cos^2 \theta_{12} \cos^2 \theta_{13} \sin^2 \theta_{14}$ (the code’s actual rotation ordering) instead of the naive $\sin^2 \theta_{14}$ already captures essentially all of the departure from the 2-flavour formula. The ν_μ channel is qualitatively different: even in the idealised limit, forcing $\theta_{14} = 0$ leaves a residual of 1.4×10^{-3} (a pure mass-hierarchy effect, not related to R_{14}/R_{24} non-commutation), which grows to 1.07×10^{-2} once both angles are simultaneously active; at realistic mixing the observed deviation (0.0821) is *larger* than what the naive-versus-exact $|U_{\mu 4}|^2$ convention alone predicts (0.0563), confirming [Chapter 1](#)’s statement that $|U_{\mu 4}|^2 = \sin^2 \theta_{24}$ is never exact once $\theta_{14} \neq 0$ – there is a genuine residual active-sector structure in the ν_μ channel that the 2-state reduction does not capture, beyond the diagonal mixing-fraction correction alone.

The notebook produces three figures: Fig 3, ν_e/ν_μ survival vs. L (log scale) for the idealised limit, numerical vs. 2-flavour formula; Fig 4, the same comparison at realistic mixing (Giunti-2017 for ν_e , IceCube for ν_μ); and Fig 5, three (E, L) heatmaps of $P(\nu_e \rightarrow \nu_e)$ for $\Delta m_{41}^2 = 0.3, 1.7, 7.0 \text{ eV}^2$ with markers for the Reactor, Gallium (BEST), LSND, MiniBooNE, and IceCube experimental (E, L) points.

4.3 `sterile3_analysis.ipynb`: experimental anomalies and the global tension

Implementation in `tpeanuts`

This notebook connects the three realistic presets to the actual experimental data motivating them ([Chapter 1](#)): reactor and Gallium (ν_e disappearance), LSND/MiniBooNE ($\nu_\mu \rightarrow \nu_e$ appearance), and IceCube (ν_μ disappearance), all evaluated with the same vacuum-propagation helper as the previous two notebooks (no Earth matter effects), and

quantifies the tension between the appearance and disappearance data sets within the minimal 3+1 model.

Result

Confirmed, directly printed output – configuration cell, the effective $e\text{--}\mu$ appearance amplitude $\sin^2(2\theta_{e\mu})_{\text{eff}} \equiv 4|U_{e4}|^2|U_{\mu4}|^2$ for each preset:

Preset	θ_{14}	θ_{24}	Δm_{41}^2	$\sin^2(2\theta_{e\mu})_{\text{eff}}$
<code>sterile_3p1_null_mixing</code>	0.00°	0.00°	1.00 eV ²	0.00000
<code>sterile_3p1_bestfit_giunti2017</code>	8.50°	7.50°	1.70 eV ²	0.00146
<code>sterile_3p1_benchmark_icecube</code>	0.00°	9.22°	0.30 eV ²	0.00000

Confirmed, directly printed output – final summary cell, the notebook's global-tension argument:

$$\text{Reactor: } \sin^2(2\theta_{14}) < 0.10 \Rightarrow \theta_{14} < 9.2^\circ$$

$$\text{IceCube: } \sin^2(2\theta_{24}) < 0.10 \Rightarrow \theta_{24} < 9.2^\circ$$

$$\text{Combined: } A_{e\mu}^{\text{eff}} < 0.10 \times \sin^2(9.2^\circ) \sim 0.10 \times 0.026 \sim 0.003$$

$$\text{LSND requires } A_{e\mu}^{\text{eff}} \sim 0.003\text{--}0.030 \Rightarrow \text{global 3+1 tension} > 3\sigma$$

Approximations and Limitations

Partially unconfirmed numbers. The notebook's Section 8 markdown summary table additionally quotes: reactor $\langle P \rangle$ at $L = 15$ m (SM 1.00, Giunti-2017 ≈ 0.97 , measured ≈ 0.94); BEST Inner survival at $L = 0.6$ m (SM 1.00, Giunti-2017 ≈ 0.94 , measured 0.79 ± 0.05); LSND peak appearance probability over $E = 20\text{--}52$ MeV (SM 0, Giunti-2017 $\approx 1.2 \times 10^{-4}$, measured ~ 0.003); IceCube survival at 1 TeV (SM ≈ 1.00 , Giunti-2017 ≈ 0.993 , a 0.7% deficit, measured deficit $< 10\%$). These numbers appear in a markdown cell that itself summarises earlier plotting cells (`st2-reactor`, `st2-gallium`, `st2-lsnd-miniboone`, `st2-icecube`); those individual cells were confirmed to have *some* recorded output but their exact raw printed text was not independently re-extracted for this document. They are reported here as notebook-stated, not as independently re-verified raw cell-output text.

The notebook produces five figures: Fig 3, reactor $\bar{\nu}_e$ survival vs. E at ILL/Bugey-3 baselines plus average rate deficit vs. L with the measured deficit band; Fig 4, ν_e survival vs. L at the ^{51}Cr source energy with BEST Inner/Outer data points and an illustrative $\theta_{14} = 20^\circ$ curve; Fig 5, $P(\bar{\nu}_\mu \rightarrow \bar{\nu}_e)$ (LSND) and $P(\nu_\mu \rightarrow \nu_e)$ (MiniBooNE) vs. E with required-signal reference lines; Fig 6, $P(\nu_\mu \rightarrow \nu_\mu)$ vs. E at Earth-diameter baseline with a θ_{24} -scan panel; and Fig 7, three $(\theta_{14}, \theta_{24})$ contour maps (appearance amplitude, reactor rate ratio, IceCube rate) with the Giunti-2017 and IceCube benchmark points marked. The notebook's own stated conclusion: the minimal 3+1 model cannot simultaneously fit appearance (LSND/MiniBooNE) and disappearance (reactor+IceCube) data – "incompatible with a combined LSND + MiniBooNE interpretation at $> 3\sigma$ " – motivating "either more complex BSM models or systematic re-examination of the anomalies," a standard conclusion in the sterile neutrino literature (see [Chapter 1](#)'s discussion of the appearance/disappearance tension).

4.4 `intrinsic3_BSM_extension_sterile.ipynb`: cross-method validation

Approximations and Limitations

This notebook's saved file contains essentially no recorded numeric results. Unlike every other notebook in this chapter, a cell-by-cell inspection of `execution_count` and `outputs` shows that every substantive comparison cell – the Section 3 SM-limit study, Section 4 solar-surface comparison, Section 5 pure-flavour Earth comparison's numeric metrics, Section 6 solar-to-detector, Section 7.1/7.2 atmosphere comparisons, and the Section 8 final summary (which contains `assert np.isfinite(...)` and a "All four physical comparisons completed successfully." print) – has **no recorded output** in the file as saved. The only two pieces of genuinely recorded evidence are a device/path configuration print (confirming the notebook was run on `cuda:0`, `torch.float64`) and a single `RuntimeWarning` raised during the Section 5 Earth comparison:

```
earth_probability_state_numerical: the 3+1 sterile extension
is active but no neutron-density data is available for this
profile/model, so the neutral-current matter term cannot be
included. Falling back to the CC-only sterile Hamiltonian
(include_matter_nc=False).
```

This confirms two things directly, and nothing else: (i) the notebook does reach the Earth-comparison cell and does exercise `earth_probability_state_numerical` with a sterile preset (`sterile_3p1_bestfit_giunti2017`) on a profile that was *not* built with neutron-density coefficients, so `resolve_include_matter_nc` (Chapter 2) correctly falls back to the CC-only sterile Hamiltonian with an explicit warning rather than silently doing so; and (ii) no comparison metric, error value, or pass/fail outcome for any of this notebook's five end-to-end scenarios (SM limit, solar surface, Earth, solar-to-detector, atmosphere) is recoverable from the saved file. The markdown "Expected results" text throughout the notebook (e.g. "active probabilities should agree to floating-point precision," "every metric must be finite") is aspirational prose, not a corroborated result, in this copy of the file.

Implementation in `tpeanuts`

The notebook's stated design (from its markdown, not from any recorded output) is a numerical-vs-perturbative cross-check with `ABSOLUTE_THRESHOLD = 1e-3` and a $|P_{\text{reference}}| \geq 10^{-5}$ floor for relative-error comparisons – structurally identical to the companion `BSM_NSI` document's `intrinsic2_BSM_extension_NSI.ipynb`, which *did* have two of its five comparisons' output recorded (see that document's Chapter 4). Re-running this notebook end to end, with a neutron-density-enabled Earth profile so that the sterile neutral-current term is actually exercised rather than silently skipped, would be a natural, low-effort follow-up to this document.

References

- [1] A. Aguilar et al. “Evidence for neutrino oscillations from the observation of anti-neutrino(electron) appearance in a anti-neutrino(muon) beam”. In: *Phys. Rev. D* 64 (2001), p. 112007. eprint: [hep-ex/0104049](#).
- [2] A. A. Aguilar-Arevalo et al. “Significant Excess of Electronlike Events in the MiniBooNE Short-Baseline Neutrino Experiment”. In: *Phys. Rev. Lett.* 121 (2018), p. 221801. eprint: [1805.12028](#).
- [3] I. Esteban, M. C. Gonzalez-Garcia, M. Maltoni, T. Schwetz, and A. Zhou. “The fate of hints: updated global analysis of three-flavor neutrino oscillations”. In: *JHEP* 09 (2020), p. 178. eprint: [2007.14792](#).
- [4] S. Gariazzo, C. Giunti, M. Laveder, and Y. F. Li. “Updated Global 3+1 Analysis of Short-BaseLine Neutrino Oscillations”. In: *JHEP* 06 (2017), p. 135. eprint: [1703.00860](#).
- [5] C. Giunti and M. Laveder. “Statistical Significance of the Gallium Anomaly”. In: *Phys. Rev. C* 83 (2011), p. 065504. eprint: [1006.3244](#).
- [6] IceCube Collaboration. “Searching for eV-scale sterile neutrinos with eight years of atmospheric neutrinos at the IceCube neutrino telescope”. In: *Phys. Rev. Lett.* 125 (2020), p. 141801. eprint: [2005.12943](#).
- [7] J. Kopp, P. A. N. Machado, M. Maltoni, and T. Schwetz. “Sterile Neutrino Oscillations: The Global Picture”. In: *JHEP* 05 (2013), p. 050. eprint: [1303.3011](#).
- [8] G. Mention et al. “The Reactor Antineutrino Anomaly”. In: *Phys. Rev. D* 83 (2011), p. 073006. eprint: [1101.2755](#).
- [9] NuFIT Collaboration. *NuFIT 5.2 (2022)*. <http://www.nu-fit.org>. 2022.
- [10] L. Wolfenstein. “Neutrino oscillations in matter”. In: *Phys. Rev. D* 17 (1978), pp. 2369–2374.