

Real Solar θ_{12} /NSI Inference at Borexino with `tpeanuts`

Fast Simulation of Neutrino Oscillations in Matter

MSc Thesis (Trabajo Fin de Máster) — Inference Companion Document

Juan Ramon Diaz Santos
`diazjuan@alumni.uv.es`

Supervisors: Roberto Ruiz de Austri Bazan & Michele Lucente

September 2026

This document accompanies `notebooks/inference/inference1_borexino.ipynb` and `notebooks/inference/inference2_borexino_nsi.ipynb`, and the `tpeanuts.detector.borexino/tpeanuts.medium.solar/tpeanuts.source.solar/tpeanuts.inference` code packages: it documents, end to end, the real Borexino solar-neutrino experiment, the official Nature 2018 open-data release it is fitted against, the solar-production, MSW-propagation, and detector-response physics implemented in `tpeanuts`, the inference machinery (likelihoods, fits, likelihood scans), and the resulting real θ_{12} measurement and Non-Standard-Interaction (NSI) LMA-Dark degeneracy map, compared against NuFIT. Blue boxes (**Theoretical Foundation**) contain the physics; green boxes (**Implementation**) link that physics to actual `tpeanuts` code; orange boxes collect approximations and limitations; amber boxes (**Result**) report a concrete numerical outcome of a notebook. Text highlighted in **yellow** marks key results/equations.

Contents

1	Introduction to the Borexino Experiment	1
1.1	Purpose and how to read this document	1
1.2	Overview	1
1.3	The neutrino-electron elastic-scattering detection principle	2
1.4	The pp-chain and the four real measured channels	3
1.5	Physics goal: the vacuum-to-matter (MSW) transition	3
2	The Official Data Release: Tables and Publications	5
2.1	Publications	5
2.2	Data-release structure and format conventions	6
2.3	Table catalogue	6
3	Solar-Neutrino and Detector Implementation in tpeanuts	8
3.1	Solar-neutrino production	8
3.2	Propagation through the Sun: adiabatic MSW conversion	9
3.3	Non-Standard Interactions: a diagonal matter-potential coupling	11
3.4	Neutrino-electron elastic-scattering cross section	12
3.5	Borexino’s target, response, background, and composed event rate	12
4	The Inference Process	16
4.1	Point-fit models: predicting P_{ee} at fixed data points	16
4.2	Likelihood statistic: asymmetric chi-square	17
4.3	Gradient-based fit procedure	17
4.4	Goodness of fit	18
4.5	Event-count model: BorexinoEventRateModel	19
4.6	2-D likelihood scan and the LMA-Dark degeneracy	19
5	Results	21
5.1	Pre-fit baseline: NuFIT versus Borexino	21
5.2	One-parameter θ_{12} fit to real Borexino data	21
5.3	Goodness of fit	22
5.4	Two-parameter fit: θ_{12} and Δm_{21}^2	23
5.5	Real event-rate shape comparison	23
5.6	NSI: the LMA-Dark degeneracy and its uncertainty map	24
5.7	Summary of inference results and comparison with NuFIT	26

Introduction to the Borexino Experiment

1.1 Purpose and how to read this document

This document is the inference-side companion to the *Theoretical Foundations of tpeanuts* document: where that document develops the generic neutrino-propagation physics implemented in `core/` and `medium/`, this one documents a single, complete, real-data application built on top of it – fitting the solar mixing angle θ_{12} (and, in a second extension, a diagonal Non-Standard-Interaction coupling ε_{ee}) to the real public measurements published by the Borexino Collaboration. Every number quoted here that is not explicitly attributed to a fitted `tpeanuts` result is a real, published, externally verifiable quantity (a survival-probability measurement or a recoil-spectrum data point), traceable through the References chapter; every illustrative or not-yet-verified quantity this project itself introduces (most notably the detector energy resolution, [Section 3.5](#)) is flagged as such wherever it appears, exactly as it is flagged in the source code it comes from.

The five chapters are organised as follows: this chapter introduces the experiment itself – detection principle, the pp-chain channels it measures, and physics goal – with no reference yet to `tpeanuts`; [Chapter 2](#) catalogues the real official Nature 2018 open-data release the code consumes, table by table; [Chapter 3](#) is the physics-and-mathematics core, developing every formula `tpeanuts.source.solar`, `tpeanuts.medium.solar`, and `tpeanuts.detector.borexino` implement, from solar production through MSW propagation to the detector response; [Chapter 4](#) documents the statistical inference machinery (likelihoods, fits, likelihood scans) and the model variants used, plain Standard Model and NSI; [Chapter 5](#) presents, interprets, and summarises the numerical results obtained from `inference1_borexino.ipynb` and `inference2_borexino_nsi.ipynb`, compared against NuFIT and against Borexino’s own published $P_{ee}(E)$ measurements.

1.2 Overview

Borexino is a large-volume, ultra-radiopure liquid-scintillator detector located deep underground at the Laboratori Nazionali del Gran Sasso (LNGS) in Italy, built with a single, focused physics goal: real-time, event-by-event spectroscopy of low-energy solar neutrinos via neutrino-electron elastic scattering (ES),

$$\nu + e^- \rightarrow \nu + e^-, \quad (1.1)$$

a real reaction every active-flavour neutrino undergoes (charged current for ν_e , neutral current for every flavour, [Chapter 3](#)), giving a continuous recoil-electron spectrum rather than a fixed detection threshold or a single reconstructed neutrino energy. This is what distinguishes Borexino from the earlier radiochemical solar-neutrino experiments (Homestake, GALLEX/SAGE): those integrated a capture rate above a fixed threshold over weeks or months and destroyed the event-by-event/timing information in the process, while Borexino resolves individual real scintillation events in time and (approximately) in energy, at a sub-MeV threshold set by its own intrinsic radiopurity ([Section 1.3](#)).

In 2018, the Borexino Collaboration published the first real, single-detector, comprehensive spectroscopic measurement of the entire low-energy pp-chain – pp, ${}^7\text{Be}$, pep, and ${}^8\text{B}$ solar

neutrinos – extracting a real survival probability $P_{ee}(E_\nu)$ at four representative energies spanning nearly two orders of magnitude in E_ν [2]. This document, and the **tpeanuts** code it accompanies, is built directly against that real published measurement: the Collaboration’s own official Open Data release of the underlying Nature 2018 figures ([Chapter 2](#)).

Approximations and Limitations

Scope of this document. **tpeanuts** reproduces the real target stoichiometry (pseudocumene, C_9H_{12}), the real tree-level Standard-Model ν - e elastic-scattering cross section, and the real adiabatic MSW conversion of solar neutrinos through a real (B16-AGSS09) standard-solar-model density and production profile, fit directly against Borexino’s own real, published $P_{ee}(E)$ points at their nominal central values and uncertainties. It does *not* reproduce Borexino’s detector geometry (vessel radii, fiducial-volume cut), its real background decomposition (^{14}C , ^{210}Po , ^{85}Kr , ^{210}Bi , cosmogenic isotopes, external gammas – [Chapter 3](#)), or its real hit-estimator energy response; the energy-resolution normalisation used for the qualitative event-rate comparison of [Section 5.5](#) is an illustrative, representative value for this detector class, explicitly *not* read off the Nature 2018 paper. [Chapter 5](#) reports, honestly and without adjustment, where this simplification does and does not matter for the final θ_{12} measurement.

1.3 The neutrino-electron elastic-scattering detection principle

Theoretical Foundation

Unlike the inverse-beta-decay reaction used by reactor experiments (a charged-current-only interaction with a fixed $\bar{\nu}_e$ threshold and a sharp delayed-coincidence tag), elastic scattering (1.1) proceeds via both the charged current (electron flavour only, exchange of a $W^\pm$) and the neutral current (every active flavour, exchange of a Z^0), giving every solar neutrino flavour – surviving ν_e and oscillated ν_μ/ν_τ alike – a real, nonzero cross section onto the same target electron ([Chapter 3](#)). This has two direct consequences for what an ES-based detector like Borexino can and cannot measure:

- **Continuous, not monoenergetic, recoil spectrum.** Even a perfectly monoenergetic incident neutrino line (e.g. the real ^7Be or pep lines) produces a continuous recoil-electron energy spectrum, since the electron recoil energy T depends on the real scattering angle as well as E_ν ([Chapter 3](#)). A single measured recoil spectrum is therefore a superposition of every solar source’s own real ES spectrum, weighted by its real cross section and real incident flux – extracting a per-source rate (as Borexino’s own analysis does, and as [Chapter 5](#) does for the fitted θ_{12}) requires disentangling these overlapping shapes, not reading off a peak position.
- **No event-by-event E_ν or direction reconstruction.** A liquid-scintillator detector measures only the real scintillation light (and, from it, the deposited energy T) with no directional or calorimetric handle on the incident neutrino itself – unlike a water Cherenkov detector (SNO, Super-Kamiokande), which reconstructs the recoil electron’s real direction and hence retains some pointing information back to the Sun. Borexino’s real physics case rests instead on its exceptional radiopurity, which pushes the practical analysis threshold low enough (~ 150 – 200 keV) to resolve the pp and ^7Be spectra directly, a regime no water Cherenkov detector can reach.

Approximations and Limitations

The real coincidence-free, single-pulse nature of an ES event also means Borexino’s real background budget is fundamentally different from a reactor or IBD-based experiment’s: without a delayed-coincidence tag, every radioactive contaminant in or near the scintillator that deposits energy in the same window as a genuine ES recoil is an irreducible background, only suppressed by extreme material radiopurity and (for alpha-emitting contaminants) pulse-shape discrimination. `tpeanuts` does not model any of these real backgrounds (Chapter 3); this is the single largest simplification this document’s real event-rate comparison (Section 5.5) needs to be read against.

1.4 The pp-chain and the four real measured channels

Theoretical Foundation

The Sun’s energy is overwhelmingly produced by the proton-proton (pp) fusion chain, a sequence of nuclear reactions each producing a real, distinct electron-neutrino source, with its own real production spectrum (continuous or a discrete line, Chapter 3) and its own characteristic energy scale. The Nature 2018 measurement [2] extracted a real survival probability at one representative energy per source, for the four channels Table 1.1 lists (also `data/detector/borexino/probability/nature2018_pee_points.csv`, Chapter 2) – together spanning the transition this document’s inference targets directly (Section 1.5).

Table 1.1: The four real pp-chain channels Borexino’s Nature 2018 $P_{ee}(E)$ measurement covers.

Source	Rep. E_ν [MeV]	Real production reaction
pp	0.267	$p + p \rightarrow {}^2\text{H} + e^+ + \nu_e$ (continuous spectrum, the dominant pp-chain branch).
${}^7\text{Be}$	0.862	${}^7\text{Be} + e^- \rightarrow {}^7\text{Li} + \nu_e$ (electron capture; a real quasi-monoenergetic line).
pep	1.440	$p + e^- + p \rightarrow {}^2\text{H} + \nu_e$ (electron-capture-assisted pp fusion; a real monoenergetic line).
${}^8\text{B}$	7.400	${}^8\text{B} \rightarrow {}^8\text{Be}^* + e^+ + \nu_e$ (continuous spectrum, the same real source SNO/Super-Kamiokande measure).

1.5 Physics goal: the vacuum-to-matter (MSW) transition

Theoretical Foundation

The real survival probability $P_{ee}(E_\nu)$ is not constant across Table 1.1’s energy range: at low energy (pp, ${}^7\text{Be}$, pep), solar neutrinos propagate essentially as in vacuum, giving $P_{ee} \approx 1 - \frac{1}{2} \sin^2(2\theta_{12})$; at high energy (${}^8\text{B}$), the Sun’s real electron density is large enough, relative to the neutrino’s energy and Δm_{21}^2 , to place propagation in the adiabatic Mikheyev-Smirnov-Wolfenstein (MSW) matter-resonance regime [6, 10], giving instead

$P_{ee} \approx \sin^2 \theta_{12}$ – a real, energy-dependent *upturn* (technically a downturn towards lower P_{ee}) between the two regimes, governed entirely by θ_{12} once Δm_{21}^2 and the Sun’s real density profile are fixed ([Chapter 3](#)). Prior experiments each accessed only one side of this transition (radiochemical experiments and KamLAND at low energy, SNO/Super-Kamiokande at high energy, each with a different detector technology and systematic budget); Borexino’s real achievement is resolving both regimes with the same detector and the same systematic framework, in one measurement [2]. This is precisely the shape [Chapter 5](#)’ one-parameter θ_{12} fit exploits, and it is why a solar-only fit has essentially no leverage on Δm_{21}^2 itself ([Section 5.4](#)): the MSW transition’s energy scale depends on Δm_{21}^2 only through where the vacuum-to-matter crossover sits, not through an oscillatory L/E pattern the way a terrestrial baseline experiment’s would.

Approximations and Limitations

This document’s second notebook, `inference2_borexino_nsi.ipynb` ([Section 5.6](#)), asks a complementary question: does the real data actually *require* the Standard-Model MSW picture above, or does an exotic Non-Standard-Interaction (NSI) coupling that flips the effective sign of the matter potential fit the same real points equally well? This is the real, known LMA-Dark degeneracy [4, 7], recovered directly from Borexino’s own data in [Section 5.6](#) rather than assumed.

The Official Data Release: Tables and Publications

2.1 Publications

Key References

- **Main measurement and data release.** M. Agostini et al. (Borexino Collaboration), *Comprehensive Measurement of pp-Chain Solar Neutrinos*, Nature **562**, 505 (2018) [2]. The official Open Data release (three plain-text tables, [Section 2.2](#)) is fetched from `borex.lngs.infn.it/wp-content/uploads/OpenData/` by `notebooks/external/solar/borexino/Borexino2_generator.ipynb` into `data/detector/borexino/`, in the tidy .csv form `tpeanuts.detector.borexino.io` consumes.
- **LMA-Dark degeneracy, original identification.** O. G. Miranda, M. A. Tortola, J. W. F. Valle, *Are Solar Neutrino Oscillations Robust?*, JHEP **10**, 008 (2006), hep-ph/0406280 [7]. First identification of the discrete $(\theta_{12}, \varepsilon_{ee})$ degeneracy [Section 5.6](#) recovers directly from the real Borexino points.
- **LMA-Dark degeneracy, updated global analysis.** I. Esteban, M. C. Gonzalez-Garcia, M. Maltoni, I. Martinez-Soler, T. Schwetz, *Updated constraints on non-standard interactions from global analysis of oscillation data*, JHEP **08**, 180 (2018), arXiv:1805.04530 [4]. Source of the reference degeneracy pair (`nsi_lma_dark_esteban2018`, `_LMA_DARK_NUFIT52_NO`) used to annotate [Chapter 5](#)’ NSI figures.
- **External comparison: NuFIT.** NuFIT Collaboration, *NuFIT 6.1 (2025)*, www.nu-fit.org [3] – website update of I. Esteban et al., JHEP **12**, 216 (2024). Source of the `_SM_NUFIT61_NO` preset used as a zero-free-parameter comparison point throughout [Chapter 5](#), never as a fit starting point or prior.
- **Standard solar model.** N. Vinyoles et al., *A New Generation of Standard Solar Models*, Astrophys. J. **835**, 202 (2017), arXiv:1611.09867 [9] – the “B16” (AGSS09 variant) model underlying the default solar production/flux/ density tables `tpeanuts.source.solar/tpeanuts.medium.solar` load, independent of Borexino’s own real data above.

2.2 Data-release structure and format conventions

Theoretical Foundation

The official Nature 2018 Open Data release ships three plain-text tables, each a direct digitisation of one figure of [2], with no format-specific reader beyond whitespace-delimited numeric parsing:

- `Nature2018_Fig2a_DATA.txt` – the real low-energy reconstructed-spectrum table, 7 whitespace-separated columns per row (`bin`, `N_h`, `energy_keV`, `width_keV`, `rate`, `rate_error`, `residual`).
- `Nature2018_Fig2b_DATA.txt` – the real high-energy radial distribution table, 5 columns (`bin`, `radius_m`, `counts`, `error`, `residual`).
- `Nature2018_Fig3_DATApoints.txt` – the real $P_{ee}(E)$ survival-probability points, laid out as four per-source text blocks (`pp`, `Be7`, `pep`, `B8`), each header followed by one or more `energy_MeV` `pee_value` `err_up` `err_down` rows.

`Borexino2_generator.ipynb` parses each table with a small, format-specific reader (whitespace tokenisation for the first two; block-header state tracking, keyed on the literal source labels `pp`, `Be7`, `pep`, `B8`, for the third – with `Be7/B8` renamed to `7Be/8B` to match the source keys used elsewhere in `tpeanuts`), and writes the tidy result to `data/detector/borexino/{observation,probability,metadata}/`. A `metadata/source.json` file records the release’s DOI, official URL, and that the probability table is *experimental points*, not a solar-model prediction, alongside every converted table.

2.3 Table catalogue

Tables 2.1 and 2.2 list every real table this project fetches and the `tpeanuts` loader that reads it; the physics each table is compared against is developed in full in Chapter 3.

Table 2.1: Real Borexino detector-observation tables.

Cached file	Real contents	Loader
<code>observation/nature2018_low_energy_spectrum.csv</code>	Real reconstructed low-energy recoil spectrum, 859 fine bins, ~ 2.4 keV wide, energy range ≈ 0.212 – 2.629 MeV; differential in the hit estimator N_h , not directly in energy (Section 3.5).	<code>load_low</code>
<code>observation/nature2018_high_energy_radial.csv</code>	Real radial distribution of high-energy events, 49 radius bins (0.05– ~ 4.85 m), summing to 3191 real counts across the sample; fetched by the generator notebook but <i>not currently consumed</i> by any <code>tpeanuts.detector.borexino</code> loader.	— (unused)

Table 2.2: Real Borexino survival-probability table (the primary fit target of [Chapter 5](#)).

Cached file	Real contents	Loader
probability/nature2018_pee_points.csv	Real measured $P_{ee}(E)$ at one representative energy per pp-chain source (pp, ${}^7\text{Be}$, pep, ${}^8\text{B}$, Table 1.1), each with a real, independently reported asymmetric 1σ uncertainty.	read directly via pa

Approximations and Limitations

Not fetched/reproduced: any real correlated systematic covariance between the four $P_{ee}(E)$ points or between recoil-spectrum bins; the real background-decomposition fit (per-species rates and pulse-shape discrimination efficiencies) that the official analysis performs to extract [Table 2.1](#)’s spectrum into per-source rates in the first place; and the real high-energy radial table ([Table 2.1](#)), whose real fiducial-volume role in the official analysis is not reproduced here since no `tpeanuts` loader currently consumes it. [Chapter 5](#) treats [Section 5.5](#)’s comparison against the real low-energy spectrum as a qualitative shape check for exactly this reason.

Solar-Neutrino and Detector Implementation in tpeanuts

This chapter develops, in full mathematical detail, every model this document’s inference is built on: where the Sun produces each pp-chain neutrino and with what spectrum (Section 3.1), how it propagates through the Sun’s real matter density to the surface (Section 3.2), how a diagonal Non-Standard-Interaction coupling modifies that propagation (Section 3.3), the real tree-level neutrino-electron elastic-scattering cross section (Section 3.4), and Borexino’s own target, response, background, and composed event-rate (Section 3.5).

3.1 Solar-neutrino production

source/solar/model.py

SolarNeutrinoSource: *real per-source production-radius distributions, total fluxes, and production spectra (continuous or discrete line), independent of how the resulting neutrinos then propagate.*

Theoretical Foundation

Each pp-chain source (Table 1.1) is produced across an extended, source-dependent real radial shell inside the Sun, not at a single point, with a real normalised production-probability density $g_s(r)$; a production-weighted average over $g_s(r)$ is what converts a radius-dependent quantity (the matter density $n_e(r)$, and hence the matter-mixing angles of Section 3.2) into a single per-source number. Each source’s real energy content is represented one of two ways:

- **Continuous spectrum** (pp, ^8B , and the others not covered by Table 1.1): a real normalised density $g(E)$ in MeV^{-1} , $\int g(E) dE = 1$, tabulated and linearly interpolated.
- **Discrete line spectrum** (^7Be , pep): one or more real fixed energies with real normalised dimensionless weights summing to 1 (e.g. ^7Be ’s two real capture lines, ground- and excited-state, at their real branching fractions) – no continuous quadrature is meaningful for a line source (Section 3.5).

The real total (flavour-summed, pre-oscillation) flux normalisation Φ_{tot} is a separate, per-source scalar multiplying either spectrum representation.

Implementation in tpeanuts

- **ContinuousSolarSpectrum**(energy_MeV, density_MeV_inverse): validates and renormalises a tabulated $g(E)$ to unit integral at construction time; `.evaluate(E_MeV)` linearly interpolates it (0 outside the tabulated range).
- **SolarLineSpectrum**(energy_MeV, weights): a real fixed set of (E_k, w_k) pairs, $\sum_k w_k = 1$.

- **SolarNeutrinoSource:** holds `production_radius`, `production_distribution(sources)` ($g_s(r)$ per source), `total_flux(source)`, `spectrum_table(source)` (returns either spectrum type above), and `mass_weights_integrate(weights_r, fractions, ...)` (the real production-weighted radial average used by `medium.solar.probability`, [Section 3.2](#)).
- `build_solar_source(None, context=...)` / `SolarNeutrinoSource.default(...)`: load the default real B16 (AGSS09) production/flux tables and legacy per-source spectrum tables [9] (provider tags "zenodo" for production/flux, "legacy" for the individual spectra, as printed by both inference notebooks' `SolarNeutrinoSource` repr).

3.2 Propagation through the Sun: adiabatic MSW conversion

`medium/solar/matter_mixing.py`, `medium/solar/adiabatic.py`, `medium/solar/probability.py`

The real charged-current MSW potential, the closed-form matter-mixing angles it induces, and the single validated entry point (`solar_probability_state`) every model in this document calls to turn oscillation parameters into a real $P_{ee}(E)$.

Theoretical Foundation

The charged-current matter potential

A solar ν_e propagating through the Sun's real electron density $n_e(r)$ picks up an extra coherent forward-scattering potential relative to ν_μ/ν_τ [10],

$$V_{CC} = \sqrt{2} G_F n_e, \quad (3.1)$$

which this project non-dimensionalises against the vacuum kinetic term $K = \Delta m^2/(2E_\nu)$ into a single ratio,

$$V_k(\Delta m^2, E_\nu, n_e) = \frac{V_{CC}}{K} = \frac{2E_\nu \hbar c \sqrt{2} G_F n_e}{\Delta m^2}, \quad (3.2)$$

sign-flipped for antineutrinos. $V_k \ll 1$ recovers the vacuum limit; $V_k \sim 1$ is the real MSW resonance region.

Closed-form matter-mixing angles

Under the adiabatic approximation (the neutrino stays in the local instantaneous matter-Hamiltonian eigenstate as $n_e(r)$ varies along its real trajectory – an excellent approximation for the Sun's real, smoothly varying density profile away from a very narrow window right at a resonance crossing), the plain 3-flavour Standard Model reduces exactly to two sequential 2-level problems, with closed-form matter-modified mixing angles [6],

$$\theta_{13}^M = \frac{1}{2} \arccos \left[\frac{\cos(2\theta_{13}) - V_k(\Delta m_{ee}^2)}{\sqrt{(\cos(2\theta_{13}) - V_k(\Delta m_{ee}^2))^2 + \sin^2(2\theta_{13})}} \right] \quad (3.3)$$

$$\theta_{12}^M = \frac{1}{2} \arccos \left[\frac{\cos(2\theta_{12}) - V'_k}{\sqrt{(\cos(2\theta_{12}) - V'_k)^2 + \sin^2(2\theta_{12}) \cos^2(\theta_{13}^M - \theta_{13})}} \right] \quad (3.4)$$

with $\Delta m_{ee}^2 = \cos^2 \theta_{12} \Delta m_{31}^2 + \sin^2 \theta_{12} \Delta m_{32}^2$ the effective atmospheric splitting and $V'_k = V_k(\Delta m_{21}^2) \cos^2 \theta_{13}^M + (\Delta m_{ee}^2 / \Delta m_{21}^2) \sin^2(\theta_{13}^M - \theta_{13})$ absorbing the small θ_{13} -sector back-reaction on the 1-2 problem. At standard solar-neutrino energies the density required for a real θ_{13} resonance exceeds the Sun's real maximum central density, so only the 1-2 (electron-density-driven) resonance is ever physically crossed.

Mass-basis production weights and the final $P_{ee}(E)$

The real production-point mass-basis weights follow directly,

$$w_1 = \cos^2 \theta_{13}^M \cos^2 \theta_{12}^M, \quad w_2 = \cos^2 \theta_{13}^M \sin^2 \theta_{12}^M, \quad w_3 = \sin^2 \theta_{13}^M, \quad (3.5)$$

production-radius-averaged over $g_s(r)$ (Section 3.1), then projected incoherently (the coherence length over an astronomical unit is negligible) onto the vacuum flavour basis via the vacuum PMNS matrix U ,

$$P_{ee}(E_\nu) = \sum_{i=1}^3 |U_{ei}|^2 \langle w_i \rangle_{g_s} = \cos^4 \theta_{13} \left[\cos^2 \theta_{12} \langle w_1 \rangle + \sin^2 \theta_{12} \langle w_2 \rangle \right] + \sin^4 \theta_{13} \langle w_3 \rangle, \quad (3.6)$$

the same real formula quoted directly in `inference1_borexino.ipynb`'s theoretical-framework section, with θ_{23}/δ_{13} dropping out exactly (they act only on the μ/τ rows of U , which (3.6)'s $|U_{ei}|^2$ projection never touches) and held fixed at an arbitrary preset value throughout this document.

Implementation in `tpeanuts`

- `matter_mixing.Vk(Deltam2, E, ne, antinu=False)` implements (3.2).
- `matter_mixing.th13_M/th12_M` implement (3.3) and (3.4) directly (the `arccos` argument is clamped to $[-1, 1]$ before the inverse trigonometric call, guaranteeing the result lies in the physical $[0, \pi/2]$ range with no extra angle-wrapping needed).
- `adiabatic.mass_weights_adiabatic_approximated(oscillation, E, ne, p_lz=None)`: the closed-form path above – plain 3-flavour Standard Model only (no NSI, no sterile extension: the 2-level reduction has no such generalisation), with an optional Landau-Zener correction [8] at the 1-2 resonance via `p_lz` (available, not exercised by either notebook this document reports on).
- `adiabatic.mass_weights_adiabatic_exact(oscillation, E, ne)`: the general- N -flavour alternative – diagonalises the full flavour-basis Hamiltonian pointwise with `torch.linalg.eigh`, exact given the adiabatic assumption (no 2-level reduction), and the only path that supports NSI (Section 3.3) or the 3+1 sterile extension. Because `eigh` sorts eigenvectors by ascending eigenvalue rather than by physical mass index (the two coincide for normal ordering but not inverted ordering), the raw output is reindexed once via a fixed permutation derived by diagonalising the vacuum-limit ($n_e = 0$) Hamiltonian and matching its eigenvectors against the PMNS matrix's own columns.
- `probability.solar_probability_mass/solar_probability_state`: the single validated entry point dispatching to exactly one of three methods – "adiabatic_approximated" (default, above), "adiabatic_exact" (above), or "numerical" (full coherent, segment-discretised propagation via `medium.solar.evolutor`, imposing no adiabaticity assumption at all) – and

raising a specific error for every unsupported method/use_LZ/BSM-extension combination (e.g. NSI requested together with "adiabatic_approximated").

Approximations and Limitations

Neither notebook this document reports on exercises the Landau-Zener correction or the "numerical" coherent-propagation path: the real Borexino $P_{ee}(E)$ fit (Section 5.2) uses "adiabatic_approximated" (its default, plain-SM-only, closed-form path), while the NSI degeneracy analysis (Section 5.6) uses "adiabatic_exact" throughout, since "adiabatic_approximated"'s 2-level reduction has no NSI generalisation. Both remain within the adiabatic approximation; validating that approximation itself against the full coherent "numerical" path for Borexino's specific energy range is not part of either notebook.

3.3 Non-Standard Interactions: a diagonal matter-potential coupling

core/BSM/bsm_nsi.py (via medium/solar/adiabatic.py's exact path)

A single diagonal NSI coupling ε_{ee} , rescaling the effective matter potential seen by ν_e – the minimal extension needed to expose the real LMA-Dark degeneracy (Section 5.6).

Theoretical Foundation

Non-Standard Interactions add a matter-sector term to the propagation Hamiltonian, on top of the Standard-Model MSW potential (3.1),

$$H = H_{\text{kin}} + H_{\text{mat}}, \quad H_{\text{mat}} = V_{CC}[\text{diag}(1, 0, 0) + \varepsilon], \quad (3.7)$$

with ε a 3×3 Hermitian matrix of dimensionless couplings in full generality. This document's NSI extension frees only the single diagonal entry ε_{ee} (every other entry fixed at exactly zero), which simply rescales the effective potential the propagating ν_e component feels,

$$V_{\text{eff}} = \sqrt{2} G_F n_e (1 + \varepsilon_{ee}). \quad (3.8)$$

Because an NSI-modified H has no closed-form mixing-angle solution analogous to (3.3) and (3.4), $P_{ee}(E_\nu)$ under NSI is computed only via the numerical "adiabatic_exact" path of Section 3.2, diagonalising (3.7) pointwise with `torch.linalg.eigh`, still under the adiabatic assumption.

Implementation in tpeanuts

`SolarNSIOscillationModel.oscillation(theta)` builds a fresh `NSIConfig(eps_ee=values["eps_ee"])` alongside the plain 3-flavour `PMNS_SM/MassSpectrum_SM` on every call, and `hamiltonian_flavour` reads `oscillation.nsi.epsilon` directly into (3.7) when building H – no branch anywhere in `medium.solar.adiabatic/probability` needs to know NSI is active beyond calling the exact-diagonalisation path (Section 3.2).

3.4 Neutrino-electron elastic-scattering cross section

detector/interaction/neutrino_electron.py

Tree-level Standard Model ν - e^- elastic scattering, differential in the electron recoil kinetic energy T , with distinct effective couplings for ν_e (charged + neutral current) and ν_μ/ν_τ (neutral current only).

Theoretical Foundation

The real differential cross section, valid for every active flavour, is

$$\frac{d\sigma}{dT}(E_\nu, T) = \frac{2G_F^2 m_e}{\pi} (\hbar c)^2 \left[g_L^2 + g_R^2 \left(1 - \frac{T}{E_\nu}\right)^2 - g_L g_R \frac{m_e T}{E_\nu^2} \right], \quad (3.9)$$

zero outside the real kinematically allowed range $0 \leq T \leq T_{\max}(E_\nu) = 2E_\nu^2/(m_e + 2E_\nu)$, with flavour-dependent effective couplings [1]

$$\nu_e \text{ (CC+NC)} : g_L = \frac{1}{2} + \sin^2 \theta_W, \quad g_R = \sin^2 \theta_W; \quad \nu_\mu, \nu_\tau \text{ (NC only)} : g_L = -\frac{1}{2} + \sin^2 \theta_W, \quad g_R = \sin^2 \theta_W. \quad (3.10)$$

Radiative corrections to (3.9) (the actual subject of [1], whose contribution here is limited to the tree-level g_L, g_R values above) are not implemented.

Implementation in tpeanuts

- `dsigma_dT(E_nu_MeV, T_MeV, g_L, g_R)` implements (3.9) directly, masking to exactly 0 outside the kinematic range.
- `NUE_COUPLINGS/NUMUTAU_COUPLINGS`: the two real coupling pairs of (3.10).
- `nue_cross_section_grid/numutau_cross_section_grid` (`E_nu_grid_MeV, T_grid_MeV`): evaluate (3.9) on the outer-product (E_ν, T) grid `detector.borexino.event_rate` (Section 3.5) folds against, shape (n_E, n_T) , cm^2/MeV .

3.5 Borexino's target, response, background, and composed event rate

detector/borexino/{parameters,response,backgrounds,event_rate,io}.py

Every Borexino-specific ingredient this project implements, composed exactly once, in `event_rate.event_rate/l`

Target: pseudocumene stoichiometry

Theoretical Foundation

Borexino's scintillator solvent is pseudocumene (1,2,4-trimethylbenzene, C_9H_{12}); feeding its real textbook composition and molar mass into the shared `detector.common.target.n_electrons` formula gives the real target electron

count at a chosen reference mass,

$$N_e^{\text{ref}} = \frac{m_{\text{ref}} \times 10^6 \text{ g/ton}}{120.19 \text{ g/mol}} \times N_A \times (9 \cdot 6 + 12 \cdot 1), \quad (3.11)$$

with $Z_C = 6$, $Z_H = 1$ the real atomic numbers of carbon and hydrogen. The reference mass/exposure pair (100 t, 1 day) is chosen only to match the unit convention of the real published spectrum (Table 2.1), not Borexino's true absolute fiducial mass or total real live time – neither is needed at this normalisation and neither is quoted here, to avoid stating from memory a number this session cannot check against the primary paper. The PPO fluor ($\sim 1.5\text{--}2.0$ g/L) and quenchers dissolved in the real scintillator are neglected: a small correction relative to the solvent's own mass fraction.

Implementation in tpeanuts

```
PSEUDOCUMENE_COMPOSITION={"C":9,"H":12}, PSEUDOCUMENE_MOLAR_MASS_G_MOL=120.19,
REFERENCE_TARGET_MASS_TON=100, REFERENCE_EXPOSURE_DAYS=1;
N_TARGET_ELECTRONS implements (3.11) once at import
time via detector.common.target.n_electrons. Default
integration grids E_NU_GRID_MEV=linspace(0.15,16.0,400),
T_GRID_MEV=TPRIME_GRID_MEV=linspace(0.05,14.0,400) span the pp-chain
neutrino energy range and the corresponding recoil range, matching
inference1_borexino.ipynb's own convention; equal true/ reconstructed grids
satisfy predicted_counts's current equal-grid requirement.
```

Energy resolution

Theoretical Foundation

A liquid-scintillator detector reconstructs deposited energy from a counted number of photoelectrons $N_{\text{pe}} \propto T$; photon counting is a Poisson process, so the *relative* statistical spread scales as $1/\sqrt{N_{\text{pe}}}$, giving an absolute resolution scaling as $\sqrt{T}$ rather than linearly,

$$\sigma(T) = a\sqrt{T}, \quad (3.12)$$

the standard functional form for this detector class, with a an empirical normalisation absorbing light yield, photon collection efficiency, and photodetector quantum efficiency.

Implementation in tpeanuts

```
sigma_MeV(T_grid_MeV, resolution_a=ENERGY_RESOLUTION_A) implements (3.12);
response_matrix feeds it into the shared
detector.common.response.gaussian_response_matrix – no Borexino-specific
convolution code beyond supplying  $\sigma(T)$ .
```

Approximations and Limitations

$a = \text{ENERGY_RESOLUTION_A} = 0.05$ (i.e. $5\%/\sqrt{\text{MeV}}$) is a representative order-of-magnitude figure for this detector class, **not** a number read off [2]. Every result in

this document that is sensitive to its exact value (only [Section 5.5](#)’s event-rate shape comparison; the $P_{ee}(E)$ -point fits of [Sections 5.2](#), [5.4](#) and [5.6](#) never touch it) should be read as illustrative until it is replaced with the paper’s own quoted resolution.

Backgrounds: an explicit placeholder

Implementation in tpeanuts

`backgrounds_MeV(n_bins)` calls the shared `detector.common.background.zero_background` directly. No real background model is implemented: Borexino’s actual low-energy budget – ^{14}C beta decay (endpoint ≈ 156 keV, setting the real practical analysis threshold), ^{210}Po alpha decay (≈ 5.3 MeV visible energy, separated by pulse-shape discrimination), ^{85}Kr , ^{210}Bi , cosmogenic isotopes, and external gammas – is each its own dedicated analysis in the Collaboration’s papers, not something to approximate here without a citable model.

Composed event rate

Theoretical Foundation

Given the oscillated flavour probabilities and the total incident flux, `event_rate` splits them into the two effective channels [Section 3.4](#) needs,

$$\Phi_e(E_\nu) = P_{ee}(E_\nu) \Phi_{\text{tot}}(E_\nu), \quad \Phi_x(E_\nu) = (1 - P_{ee}(E_\nu)) \Phi_{\text{tot}}(E_\nu), \quad (3.13)$$

exact for a two-flavour-equivalent treatment of solar propagation (surviving flux is ν_e with probability P_{ee} , “everything else” – ν_μ or ν_τ , equal ES cross section – with probability $1 - P_{ee}$), and folds them into predicted counts exactly as the shared `detector.common.event_rate.predicted_counts` machinery does for every detector in this project,

$$\frac{dR}{dT}(T) = N_e^{\text{ref}} \int dE_\nu \left[\Phi_e(E_\nu) \frac{d\sigma_e}{dT} + \Phi_x(E_\nu) \frac{d\sigma_x}{dT} \right], \quad \frac{dR}{dT'}(T') = \int dT R(T'|T) \frac{dR}{dT}(T), \quad (3.14)$$

$$N_i^{\text{pred}} = \tau \int_{T'_i}^{T'_{i+1}} dT' \varepsilon(T') \frac{dR}{dT'}(T') + N_i^{\text{bkg}}, \quad (3.15)$$

with $\tau = \text{exposure_days} \times 86,400$ s the exposure, ε an optional hard step-function analysis threshold (uniformly 1 if none is given), and N_i^{bkg} currently always zero ([Section 3.5](#), backgrounds). This composed function deliberately never imports `medium.solar` or any propagation code: it takes P_{ee}/Φ_{tot} as plain, already-computed tensors, wired together upstream by `BorexinoEventRateModel` ([Section 4.1](#)).

Implementation in tpeanuts

- `event_rate(probabilities, flux_tot_MeV, bin_edges_MeV, ..., exposure_days=REFERENCE_EXPOSURE_DAYS)` implements (3.15) for a continuous-spectrum source, calling `detector.common.event_rate.predicted_counts` with the real target, response, efficiency, and background plugged in. Returns counts per bin in events/`exposure_days` days – at the default 1 day, directly a rate.

- `line_event_rate(probabilities, total_flux, line_weights, line_energy_MeV, bin_edges_MeV, ...)`: the direct parallel for monoenergetic lines (^7Be , pep), wiring through `predicted_counts_discrete` instead – skipping the neutrino-energy quadrature that is meaningless for a discrete line – otherwise identical exposure/background handling.
- Fully differentiable w.r.t. `probabilities` (and hence, through it, w.r.t. whatever upstream oscillation parameters produced it); every other input (grids, cross sections, response, target, background) is a constant.

Loading the real published spectrum

Implementation in tpeanuts

`load_low_energy_spectrum(path=None)` reads `observation/nature2018_low_energy_spectrum.csv` (Table 2.1) via `pandas.read_csv`. Its published ordinate is differential in the hit estimator N_h , not in energy; each row’s calibrated `width_keV` gives the local Jacobian $dN_h/dE = 1/dE$, applied to convert the table’s per-hit-bin `rate/rate_error` into a density per MeV,

$$\text{value} = \frac{\text{rate}}{\text{width_MeV}}, \quad \sigma = \frac{\text{rate_error}}{\text{width_MeV}}, \quad (3.16)$$

returned as a symmetric `Observation` with `bin_width_MeV` set to the real calibrated energy width; multiplying `value` back by `bin_width_MeV` recovers the table’s original published rate.

Approximations and Limitations

Read before comparing against `event_rate`’s predicted counts. Equation (3.16) is only a change of variable on the published N_h -binned rate – *not* a full Borexino $R(N_h|E)$ response model – and adds no background model either. Section 5.5 treats any comparison between the loaded `Observation` and `event_rate`’s output as a qualitative shape check only, restricted to a normalisation-free quantity (both curves rescaled to unit area), not a bin-for-bin fit.

The Inference Process

This chapter documents the statistical inference machinery shared by both notebooks this document reports on: the point-fit model contract that lets a handful of real (E, P_{ee}^{obs}) measurements drive a gradient-based fit (Section 4.1), the asymmetric-uncertainty likelihood statistic (Section 4.2), the gradient-based fit procedure and its Laplace uncertainty (Section 4.3), goodness-of-fit diagnostics (Section 4.4), the event-count forward-folding model (Section 4.5), and the 2-D likelihood scan used to map the real NSI degeneracy (Section 4.6).

4.1 Point-fit models: predicting P_{ee} at fixed data points

`inference/model_solar.py`

`SolarSMOscillationModel`, **`SolarNSIOscillationModel`**, **`SolarPointModel`**: differentiable $P_{ee}(E)$ models, adapted to the single-argument `predict(theta)` contract `fit_lbfgs` (Section 4.3) expects.

Theoretical Foundation

Plain Standard-Model path

`SolarSMOscillationModel` exposes a free-parameter vector drawn from $(\theta_{12}, \theta_{13}, \Delta m_{21}^2, \Delta m_{3l}^2)$ – every entry not in `free` stays fixed at its preset value – and builds a fresh **`OscillationParameters`** (plain 3-flavour `PMNS_SM` + `MassSpectrum_SM`, no NSI, no sterile extension) from it on every call, so autograd differentiates the predicted P_{ee} with respect to `theta` through Chapter 3’s full MSW chain. `.predict_pee(theta, medium, source, sources, energies_MeV)` evaluates (3.6) independently at each (source, energy) pair – a single-energy point evaluation, not a spectrum-averaged rate, matching exactly how Borexino’s own real $P_{ee}(E)$ points are themselves reported (one representative energy per source, Table 1.1).

NSI extension

`SolarNSIOscillationModel` adds a single further free-parameter name, ε_{ee} , on top of the same four, and always evaluates `predict_pee` via `method="adiabatic_exact"` (Section 3.2), since the closed-form path has no NSI generalisation.

Binding a model to a fixed set of data points

`SolarPointModel` binds the (fixed, per-analysis) evaluation points – `medium, source, sources, energies_MeV` – at construction time, so the wrapped oscillation model only needs to expose its own `predict_pee(theta, medium, source, sources, energies_MeV)` to satisfy the plain `predict(theta)` contract every fitting/scanning utility in this document expects.

Implementation in tpeanuts

- `SolarSMOscillationModel.from_preset(preset_name, free=...)` builds the model and a starting `theta0` leaf tensor from a named entry of `tpeanuts.config.presets.OSCILLATION_PRESETS` (e.g. `"_SM_NUFIT61_NO"`).
- `SolarNSIOscillationModel.from_preset(preset_name, free=...), eps_ee0=0.0)`: identical construction, plus an explicit ε_{ee} starting value (no oscillation preset carries NSI information).
- `SolarPointModel(oscillation_model, medium, source, sources, energies_MeV):` `.predict(theta)` delegates directly to `oscillation_model.predict_pee(theta, ...)`.

4.2 Likelihood statistic: asymmetric chi-square

`inference/likelihood.py`

`chi2_asymmetric`: the one statistic every fit in this document minimises, reused unchanged from the rest of `tpeanuts`.

Theoretical Foundation

Each of the four real $P_{ee}(E)$ measurements carries its own independently reported, generally asymmetric 1σ uncertainty (Table 1.1); comparing a prediction against it uses a one-sided-uncertainty Gaussian statistic, selecting σ_+ or σ_- by the sign of the residual,

$$-2 \ln L = \chi^2(\theta) = \sum_{i=1}^N \left(\frac{P_{ee}^{\text{pred}}(E_i; \theta) - P_{ee}^{\text{obs}}(E_i)}{\sigma_i} \right)^2, \quad \sigma_i = \begin{cases} \sigma_{+,i}, & P_{ee}^{\text{pred}} \geq P_{ee}^{\text{obs}}, \\ \sigma_{-,i}, & \text{otherwise,} \end{cases} \quad (4.1)$$

with $N = 4$ throughout Chapter 5 (one term per real source of Table 1.1).

Implementation in tpeanuts

`chi2_asymmetric(prediction, value, sigma_minus, sigma_plus)` implements (4.1) exactly, differentiable w.r.t. `prediction`; passed as the `likelihood` argument to `fit_lbfgs/loglik_grid` (Sections 4.3 and 4.6) throughout this document. `tpeanuts.inference.likelihood` also provides `poisson_nll` (the Poisson statistic used elsewhere in this project for absolute per-bin counts, e.g. Daya Bay's inference document) and `gaussian_prior_penalty`; neither is exercised here, since this document's real observations are all asymmetric-Gaussian point measurements and no Gaussian-constrained nuisance parameter is fit.

4.3 Gradient-based fit procedure

`inference/fit.py`

`fit_lbfgs/minimize_lbfgs`: L -BFGS point estimate plus Laplace/Fisher covariance, for any `DifferentiableModel` – reused unchanged from the rest of `tpeanuts`.

Theoretical Foundation

L-BFGS

L-BFGS (Limited-memory Broyden-Fletcher-Goldfarb-Shanno) is a gradient-based quasi-Newton method that builds a limited-memory approximation to the inverse Hessian from a short history of parameter/gradient displacements, rather than constructing the full Hessian [5]. Every fit in this document starts from a deliberately uninformative guess (e.g. $\theta_{12} = 45^\circ$, maximal mixing – never the NuFIT value, which is used only as an independent comparison point, [Section 5.2](#)) and uses `torch.optim.LBFGS` with the exact autograd gradient of (4.1) through every real physics step of [Chapter 3](#) (no finite differences).

Parameter rescaling

Free parameters spanning many orders of magnitude (e.g. $\theta_{12} \sim \mathcal{O}(1)$ rad against $\Delta m_{21}^2 \sim \mathcal{O}(10^{-4})$ eV² in [Section 5.4](#)’s two-parameter fit) starve L-BFGS’s identity-initialised inverse-Hessian approximation of a usable step size; a dimensionless $u = \theta/\text{scale}$ is fit instead, with scale set from the loss gradient at θ_0 and undone only in the returned $\hat{\theta}$ – every model always sees physical-unit θ .

Laplace/Fisher covariance

The marginal uncertainty on each fitted parameter is the Laplace approximation: with $\mathcal{I} = \frac{1}{2} \nabla_{\theta}^2 \chi^2$ the Fisher information at the optimum (the autograd Hessian of the total statistic),

$$\Sigma = \mathcal{I}^{-1}, \quad \sigma_j = \sqrt{\Sigma_{jj}}. \quad (4.2)$$

A direction the data barely constrains gives a Hessian eigenvalue near zero; small or negative eigenvalues (numerical round-off) are floored at a tiny fraction of the largest one before inversion, so an unconstrained direction gets a large but finite σ_j instead of nan – exactly the behaviour [Section 5.4](#) reports for Δm_{21}^2 .

Implementation in tpeanuts

`fit_lbfgs(model, theta0, value, sigma_minus, sigma_plus, likelihood="chi2_asymmetric")` returns a `FitResult` (`theta_hat`, `param_names`, `chi2_history`, and a `sigma` property implementing (4.2)); every fit in [Chapter 5](#) uses this one function.

4.4 Goodness of fit

Theoretical Foundation

With $N = 4$ real data points and k free parameters, a fit has $\nu = N - k$ degrees of freedom. Under the standard large-sample approximation, $\chi_{\min}^2$ at the true parameters is itself χ^2 -distributed with ν degrees of freedom, giving the p -value

$$p = P(\chi_{\nu}^2 \geq \chi_{\min}^2) = 1 - F_{\chi_{\nu}^2}(\chi_{\min}^2), \quad (4.3)$$

evaluated with `scipy.stats.chi2.sf` in [Section 5.3](#).

4.5 Event-count model: `BorexinoEventRateModel`

detector/borexino/inference_model.py

Wraps a `SolarSMOscillationModel`'s $P_{ee}(E)$ curve with `detector.borexino.event_rate` (Section 3.5), giving a `predict(theta)` that returns real per-bin event counts rather than P_{ee} – the model Section 5.5's shape comparison uses.

Theoretical Foundation

For each requested solar source name, `BorexinoEventRateModel` evaluates P_{ee} on the full energy grid (continuous sources) or at the source's own discrete line energies (Section 3.1), folds it through `event_rate/line_event_rate` (Eq. (3.15)) with the model's own bound `bin_edges_MeV/exposure_days`, and sums the per-source contributions before adding the (currently zero) detector background once – never per source, avoiding a background double-count when several sources are requested together.

Implementation in tpeanuts

```
BorexinoEventRateModel(oscillation_model, medium, source, source_names,
bin_edges_MeV, exposure_days=REFERENCE_EXPOSURE_DAYS):    .free    forwards
oscillation_model.free; .predict(theta) builds oscillation = oscillation_model
.oscillation(theta), then for each source_name in source_names dispatches on whether
source.spectrum_table(source_name) is a SolarLineSpectrum (calls line_event_rate)
or continuous (calls event_rate on E_nu_grid_MeV), fully differentiable w.r.t. theta
in both branches (validated directly against a hand-built line_event_rate call in
detector/borexino/test/test_inference_model.py).
```

4.6 2-D likelihood scan and the LMA-Dark degeneracy

inference/scan.py

`loglik_grid`: evaluates $-2 \ln L$ on a (θ_x, θ_y) grid; `CHI2_LEVELS_2D`: the standard Wilks' theorem confidence thresholds at 2 degrees of freedom.

Theoretical Foundation

Confidence regions from a likelihood scan

For a best-fit point $\hat{\theta}$ minimising $-2 \ln L(\theta)$, the standard joint confidence region for two of its components scans $\Delta(-2 \ln L)(\theta_x, \theta_y) = -2 \ln L(\theta_x, \theta_y) - (-2 \ln L)_{\min}$ on a grid and draws contours where it crosses a confidence-level threshold. Wilks' theorem states that, asymptotically, $\Delta(-2 \ln L)$ at the true parameters is χ^2 -distributed with 2 degrees of freedom (one per scanned parameter), giving the standard table

$$\text{CL} = 68\% : 2.30, \quad 90\% : 4.61, \quad 95\% : 6.18, \quad 99\% : 9.21. \quad (4.4)$$

The LMA-Dark degeneracy relation

At $\varepsilon_{ee} = -2$, (3.8) gives $V_{\text{eff}} \rightarrow -\sqrt{2}G_F n_e$ – the same magnitude, opposite sign. Reflecting θ_{12} into the second octant compensates the sign flip exactly in the adiabatic limit, so

$P_{ee}^{\text{adiabatic}} = \cos^2 \theta_{12}$ is unchanged either way. More generally [4, 7],

$$\theta_{12} \longrightarrow \frac{\pi}{2} - \theta_{12}, \quad \varepsilon_{ee} \longrightarrow -(2 + \varepsilon_{ee}) \quad (4.5)$$

leaves the adiabatic $P_{ee}(E)$ numerically unchanged – a real, testable prediction [Section 5.6](#) verifies directly against two independently-converged real-data fits.

Implementation in tpeanuts

- `loglik_grid(model, theta_baseline, param_x, x_grid, param_y, y_grid, value, sigma_minus, sigma_plus, likelihood="chi2_asymmetric")`: with exactly two free parameters in `model.free` (both point-fit models here, once the SM 4-parameter vector is restricted to `("theta12", "eps_ee")`), nothing remains to profile – `profile_others` is moot, and [Section 5.6](#)’s scans are genuine 2-D likelihood surfaces of the model’s full free set.
- `CHI2_LEVELS_2D`: the dictionary implementing (4.4) (Numerical Recipes Table 15.6.1 / PDG Statistics Review), passed directly to `matplotlib`’s `contour(..., levels=...)` against a `loglik_grid` output (after subtracting its own minimum).
- `calibrate_delta_threshold`: a Monte-Carlo (parametric-bootstrap) alternative to (4.4), checking whether Wilks’ theorem actually holds at a given exposure/statistics level instead of assuming it. Available in `inference.scan`, but not exercised by either notebook this document reports on – deferred, by `inference2_borexino_nsi.ipynb`’s own account, to a separate code-validation notebook.

Results

This chapter presents, interprets, and summarises the real numerical results of `notebooks/inference/inference1_borexino.ipynb` (Sections 5.1 to 5.5) and `notebooks/inference/inference2_borexino_nsi.ipynb` (Section 5.6), executed end to end against the real Borexino Nature 2018 $P_{ee}(E)$ points (Chapter 2) using the physics of Chapter 3 and the inference machinery of Chapter 4. Every number below is a real, reproducible output of one of those two notebooks, not illustrative.

5.1 Pre-fit baseline: NuFIT versus Borexino

Before running any fit, `inference1_borexino.ipynb` first evaluates how far the current NuFIT 6.1 global-fit θ_{12} sits from the real Borexino points, using the plain adiabatic model of Eq. (3.6) with θ_{12} fixed at the NuFIT value – a zero-free-parameter χ^2 , not a fit, setting expectations for whether a fit has anywhere non-trivial to move.

Result

At NuFIT 6.1's $\theta_{12} = 33.76^\circ$ [3], every real pull remains below one in magnitude and the fixed-point statistic is small (Table 5.1),

$$\chi^2(\theta_{12} = 33.76^\circ) = 0.852 \quad (4 \text{ points}, 0 \text{ free parameters}). \quad (5.1)$$

Table 5.1: Real Borexino $P_{ee}(E)$ points vs. the NuFIT 6.1 prediction, θ_{12} fixed (no fit).

Source	E [MeV]	P_{ee}^{obs}	P_{ee}^{NuFIT}	pull
pp	0.267	0.566	0.5430	−0.250
^7Be	0.862	0.532	0.5252	−0.126
pep	1.440	0.428	0.5152	+0.765
^8B	7.400	0.390	0.3508	−0.435

5.2 One-parameter θ_{12} fit to real Borexino data

Minimising (4.1) over θ_{12} alone, starting from a deliberately uninformative guess (45° , maximal mixing) rather than the NuFIT value – a fit that only works when started near the answer is not a meaningful demonstration.

Result

$$\hat{\theta}_{12} = 34.594^\circ \pm 4.150^\circ, \quad (5.2)$$

starting from 45.000° (uninformative), with χ^2 falling from 3.039 at the start to 0.809 at convergence (10 L-BFGS closure evaluations). Relative to NuFIT 6.1's 33.76° (reference

only, never used as a prior or starting point): $|\hat{\theta}_{12} - \theta_{12}^{\text{NuFIT}}|/\sigma_{\hat{\theta}_{12}} = 0.20$ – fully consistent.




Figure 5.1: The real Borexino $P_{ee}(E)$ points, the uninformative 45° starting curve, the NuFIT 6.1 reference curve, and the fitted $\theta_{12} = 34.59^\circ$ curve (all evaluated for the ^8B source, for reference across the full pp-chain energy range).

Figure 5.1 shows the real vacuum-to-MSW transition (Section 1.5) directly: the fitted curve tracks the real low-energy (vacuum-dominated) and high-energy (^8B , matter-enhanced) data simultaneously, with only θ_{12} free.

5.3 Goodness of fit

With $N = 4$ real points and $k = 1$ free parameter ($\nu = 3$ degrees of freedom):

Result

$$\chi_{\min}^2 = 0.809, \quad \chi_{\min}^2/\nu = 0.270, \quad p = 0.847, \quad (5.3)$$

a small but statistically unremarkable χ^2 with no large pull (Table 5.2) – a real, close, honest fit, not an overfit to 4 points.

Table 5.2: Real per-point pull at the one-parameter best fit ($\hat{\theta}_{12} = 34.594^\circ$).

Source	Pull at best fit
pp	−0.351
^7Be	−0.279
pep	+0.698
^8B	−0.348

5.4 Two-parameter fit: θ_{12} and Δm_{21}^2

Repeating the fit with Δm_{21}^2 also free tests directly the claim of [Section 1.5](#): with only four fixed-energy points and no genuine L/E oscillatory handle, $P_{ee}(E)$ should be only weakly sensitive to Δm_{21}^2 away from the lowest energies.

Result

$$\hat{\theta}_{12} = 38.90^\circ \pm 5.64^\circ \quad (\text{cf. } 34.59^\circ \pm 4.15^\circ \text{ one-parameter above}), \quad (5.4)$$

$$\Delta \hat{m}_{21}^2 = 0.0247 \pm 0.442 \text{ (same units as the } 10^{-4} \text{ eV}^2 \text{ starting value)} \quad \text{essentially unconstrained}, \quad (5.5)$$

starting from $(45.0^\circ, 1.0 \times 10^{-4})$, with χ^2 falling only from 3.039 to 2.756 – θ_{12} shifts by little more than 1σ once Δm_{21}^2 is also free, and Δm_{21}^2 's own fitted relative uncertainty exceeds its central value by more than an order of magnitude. Exactly as anticipated: these four fixed-energy points do not carry the L/E information a genuine Δm_{21}^2 measurement needs, motivating complementary reactor/day-night observables elsewhere in this project's inference chain.

5.5 Real event-rate shape comparison

Extending the fitted $\theta_{12} = 34.59^\circ$ from [Section 5.2](#) to a predicted recoil-electron *event-count* spectrum ([Eq. \(3.15\)](#)), via `BorexinoEventRateModel` restricted to the ^8B source (the dominant real contributor in the recoil-energy window plotted below), and comparing its shape against Borexino's own real published low-energy spectrum ([Section 3.5](#)).

Result

At $N_e^{\text{ref}} = 3.3069 \times 10^{31}$ electrons (100 t reference mass, [Eq. \(3.11\)](#)), 20 bins spanning 0.50–6.50 MeV, and an illustrative exposure of 1000 days (≈ 2.7 “detector-years” at the reference mass – *not* Borexino's real Phase-II exposure, chosen only to give an $\mathcal{O}(10^2)$ -count illustrative total), the predicted total is

$$N_{\text{pred}}^{\text{total}}(\hat{\theta}_{12} = 34.59^\circ, 1000 \text{ days}) = 258.6 \text{ counts}. \quad (5.6)$$

Approximations and Limitations

[Figure 5.2](#) shows the predicted ^8B -only shape falling far more gently than the real published spectrum, which is dominated at these energies by the real $\text{pp}/^7\text{Be}/\text{pep}$ contributions this ^8B -only slice deliberately ignores, on top of the approximate N_h -to-energy Jacobian ([Section 3.5](#)), the illustrative energy resolution ([Section 3.5](#)), and the entirely absent background model ([Section 3.5](#)). Any residual disagreement is at least as likely to come from these known, explicitly flagged simplifications as from oscillation physics – [Fig. 5.2](#) is read as a qualitative sanity check that the forward-folding chain produces a physically reasonable, positive, broadly-shaped spectrum, not as evidence for or against the fitted θ_{12} .




Figure 5.2: Predicted ${}^8\text{B}$ -only recoil shape (fitted θ_{12}) vs. the real Borexino low-energy spectrum, both rescaled to unit area over the plotted range – a qualitative shape check, not a bin-for-bin fit.

5.6 NSI: the LMA-Dark degeneracy and its uncertainty map

χ^2 landscape: two valleys

A plain $35 \times 35 = 1225$ -point grid scan of $\chi^2(\theta_{12}, \varepsilon_{ee})$ (no optimiser involved) over a broad range of both parameters, visualising the full real likelihood surface before any fit commits to one region of it.

Result

The scan resolves two separate valleys near the two reference regions (standard LMA, $\varepsilon_{ee} \approx 0$; LMA-Dark, $\varepsilon_{ee} \approx -2$), with a finite grid minimum

$$\chi_{\text{grid,min}}^2 = 0.804. \quad (5.7)$$

Precise fits from both sides

Running the gradient-based fit twice, from deliberately uninformative starting guesses on either side of $\theta_{12} = 45^\circ$ (not the reference preset values, and not each other's expected answer).

Result

LMA-side fit (start 25.00° , $+0.50$): $\hat{\theta}_{12} = 34.50^\circ \pm 4.08^\circ$, $\hat{\varepsilon}_{ee} = -0.225 \pm 0.685$, $\chi^2 : 10.714 \rightarrow 0.737$, (5.8)

Dark-side fit (start 65.00° , -1.20): $\hat{\theta}_{12} = 55.50^\circ \pm 4.08^\circ$, $\hat{\varepsilon}_{ee} = -1.782 \pm 0.700$, $\chi^2 : 18.762 \rightarrow 0.737$, (5.9)

$$|\chi_{\text{LMA}}^2 - \chi_{\text{dark}}^2| = 0.0002 \quad - \text{indistinguishable given the real data.} \quad (5.10)$$




Figure 5.3: Real $\chi^2(\theta_{12}, \varepsilon_{ee})$ against the Borexino data, showing two distinct low- χ^2 valleys around the standard-LMA and LMA-Dark reference points.

Two deliberately different, uninformative starting guesses converge to two well-separated points in $(\theta_{12}, \varepsilon_{ee})$ space with statistically indistinguishable χ^2 against the real Borexino data – the LMA-Dark degeneracy, recovered by gradient descent rather than assumed.

Testing the degeneracy relation

Comparing the dark-side fit's actual $(\theta_{12}, \varepsilon_{ee})$ against what (4.5) predicts *from the LMA-side fit alone*.

Result

Quantity	Predicted from LMA fit	Actual dark-side fit	Diff
$\theta_{12}^{\text{dark}}$ [deg]	55.50	55.50	+0.00
$\varepsilon_{ee}^{\text{dark}}$	-1.775	-1.782	+0.007

Both quantities agree with (4.5)'s prediction well within the dark-side fit's own 1σ uncertainty – the degeneracy relation is quantitatively confirmed against real data, not merely qualitatively observed.

$P_{ee}(E)$ curves: where the degeneracy holds

Plotting the full $P_{ee}(E)$ curve at both best-fit points across the full 0.15–16 MeV pp-chain range, and their absolute difference, tests whether the degeneracy (exact only in the strict adiabatic limit) also holds away from it.

Result

$$\max_E |P_{ee}^{\text{LMA}}(E) - P_{ee}^{\text{dark}}(E)| = 0.0014 \quad \text{at } E = 16.00 \text{ MeV (the sampled grid's high-energy end),} \quad (5.11)$$

vanishing at the low-energy end ($E = 0.15 \text{ MeV}$: difference = 0.0000) and peaking, still at the sub-percent level, in the vacuum-to-matter transition region around a few MeV.




Figure 5.4: Left: both best-fit $P_{ee}(E)$ curves against the real Borexino data, visually indistinguishable. Right: their absolute difference, non-monotonic – smallest at both energy extremes (exact by construction in the deep-vacuum and deep-adiabatic-MSW limits), largest, but still small, in the transition region.

Approximations and Limitations

A real resolution of LMA vs. LMA-Dark needs independent information (e.g. reactor/accelerator data, or non-oscillation NSI constraints), not more solar $P_{ee}(E)$ points alone – exactly why this is a known, real open degeneracy [4, 7], not a defect of this fit.

Calibrated uncertainty map

Mapping the 2-D $\Delta(-2 \ln L)$ surface (Eq. (4.4)) around the LMA-side real-data fit via `loglik_grid`, drawing the 68%/90%/95% Wilks'-theorem contours.

Result

On a $30 \times 30 = 900$ -point grid,

$$\chi_{\min}^2 = 0.777 \quad \text{at } \theta_{12} = 35.34^\circ, \quad \varepsilon_{ee} = -0.24, \quad (5.12)$$

close to the LMA-side fit of Section 5.6's precise fit above, with the 68%/90%/95% contours nested as expected.

5.7 Summary of inference results and comparison with NuFIT

Table 5.3 collects every real fitted value obtained in this chapter alongside the independent NuFIT 6.1 global-fit value [3], as used elsewhere in this project's own preset registry




Figure 5.5: Calibrated $\Delta(-2\ln L)$ contours (Wilks, 2 degrees of freedom) around the real-data LMA-side fit.

(tpeanuts.config.presets._SM_NUFIT61_N0).

Table 5.3: Final comparison: tpeanuts real-data inference vs. NuFIT 6.1.

Quantity	tpeanuts (Borexino, this chapter)	NuFIT 6.1 [†]
θ_{12} [deg], 1-param fit (§5.2)	34.59 ± 4.15	33.76
θ_{12} [deg], 2-param fit (§5.4)	38.90 ± 5.64	33.76
Δm_{21}^2 , 2-param fit (§5.4)	essentially unconstrained	$7.537 \times 10^{-5} \text{ eV}^2$
θ_{12} [deg], NSI LMA-side (§5.6)	34.50 ± 4.08	33.76
ε_{ee} , NSI LMA-side (§5.6)	-0.225 ± 0.685	0 (SM)

[†]NuFIT 6.1’s own global-fit best-fit central value [3], itself informed by, among other inputs, Borexino’s own real published measurement, not an independent cross-check.

Result

Final assessment. The real one-parameter θ_{12} fit (Section 5.2) recovers the solar mixing angle to $\sim 4^\circ$ precision from just four real points, fully consistent with NuFIT within 0.2σ – a meaningful closed-loop demonstration on real data, not a competitive measurement (Chapter 3’s scope notebbox: no spectral/detector-response folding, no inter-point correlations, only 4 points). The two-parameter extension confirms Δm_{21}^2 is essentially unconstrained by these four fixed-energy points, exactly as the theoretical framework predicts. The NSI extension (Section 5.6) recovers the real, known LMA-Dark degeneracy directly from the data, both qualitatively (two independently-converged, statistically indistinguishable fits) and quantitatively (the degeneracy relation (4.5)

verified to sub-degree/sub-0.01 precision), and maps a calibrated confidence region around the standard-LMA solution – a genuine open degeneracy this document surfaces rather than resolves, since resolving it needs independent, non-solar information.

Recommended next step. The event-rate shape comparison ([Section 5.5](#)) is the weakest link in this document, resting on an illustrative (not paper-verified) energy resolution and no background model; replacing `ENERGY_RESOLUTION_A` with Borexino’s own quoted resolution and folding in even a schematic background model would let a future extension of this document attempt an actual bin-by-bin spectral fit, rather than the qualitative shape check reported here.

References

- [1] J. N. Bahcall, M. Kamionkowski, and A. Sirlin. “Solar neutrinos: Radiative corrections in neutrino-electron scattering experiments”. In: *Phys. Rev. D* 51 (1995), pp. 6146–6158. eprint: [astro-ph/9502003](#).
- [2] Borexino Collaboration, M. Agostini, et al. “Comprehensive measurement of pp-chain solar neutrinos”. In: *Nature* 562 (2018), pp. 505–510. DOI: [10.1038/s41586-018-0624-y](#).
- [3] I. Esteban, M. C. Gonzalez-Garcia, M. Maltoni, I. Martinez-Soler, J. P. Pinheiro, and T. Schwetz. *NuFIT 6.1 (2025)*. 2025. URL: <http://www.nu-fit.org>.
- [4] I. Esteban, M. C. Gonzalez-Garcia, M. Maltoni, I. Martinez-Soler, and T. Schwetz. “Updated constraints on non-standard interactions from global analysis of oscillation data”. In: *JHEP* 08 (2018), p. 180. eprint: [1805.04530](#).
- [5] D. C. Liu and J. Nocedal. “On the Limited Memory BFGS Method for Large Scale Optimization”. In: *Mathematical Programming* 45 (1989), pp. 503–528. DOI: [10.1007/BF01589116](#).
- [6] S. P. Mikheyev and A. Y. Smirnov. “Resonance Amplification of Oscillations in Matter and Spectroscopy of Solar Neutrinos”. In: *Sov. J. Nucl. Phys.* 42 (1985), pp. 913–917.
- [7] O. G. Miranda, M. A. Tortola, and J. W. F. Valle. “Are Solar Neutrino Oscillations Robust?” In: *Journal of High Energy Physics* 10 (2006), p. 008. DOI: [10.1088/1126-6708/2006/10/008](#). eprint: [hep-ph/0406280](#).
- [8] S. J. Parke. “Nonadiabatic Level Crossing in Resonant Neutrino Oscillations”. In: *Phys. Rev. Lett.* 57 (1986), pp. 1275–1278.
- [9] N. Vinyoles et al. “A New Generation of Standard Solar Models”. In: *Astrophys. J.* 835 (2017), p. 202. eprint: [1611.09867](#).
- [10] L. Wolfenstein. “Neutrino oscillations in matter”. In: *Phys. Rev. D* 17 (1978), pp. 2369–2374.